

Memory Profiling in Pharo

Sebastian JORDAN MONTAÑO

sebastian.jordan@inria.fr

Inria, Univ. Lille, CNRS, Centrale Lille, UMR 9189 - CRISTAL



Université
de Lille



November 2023

About me

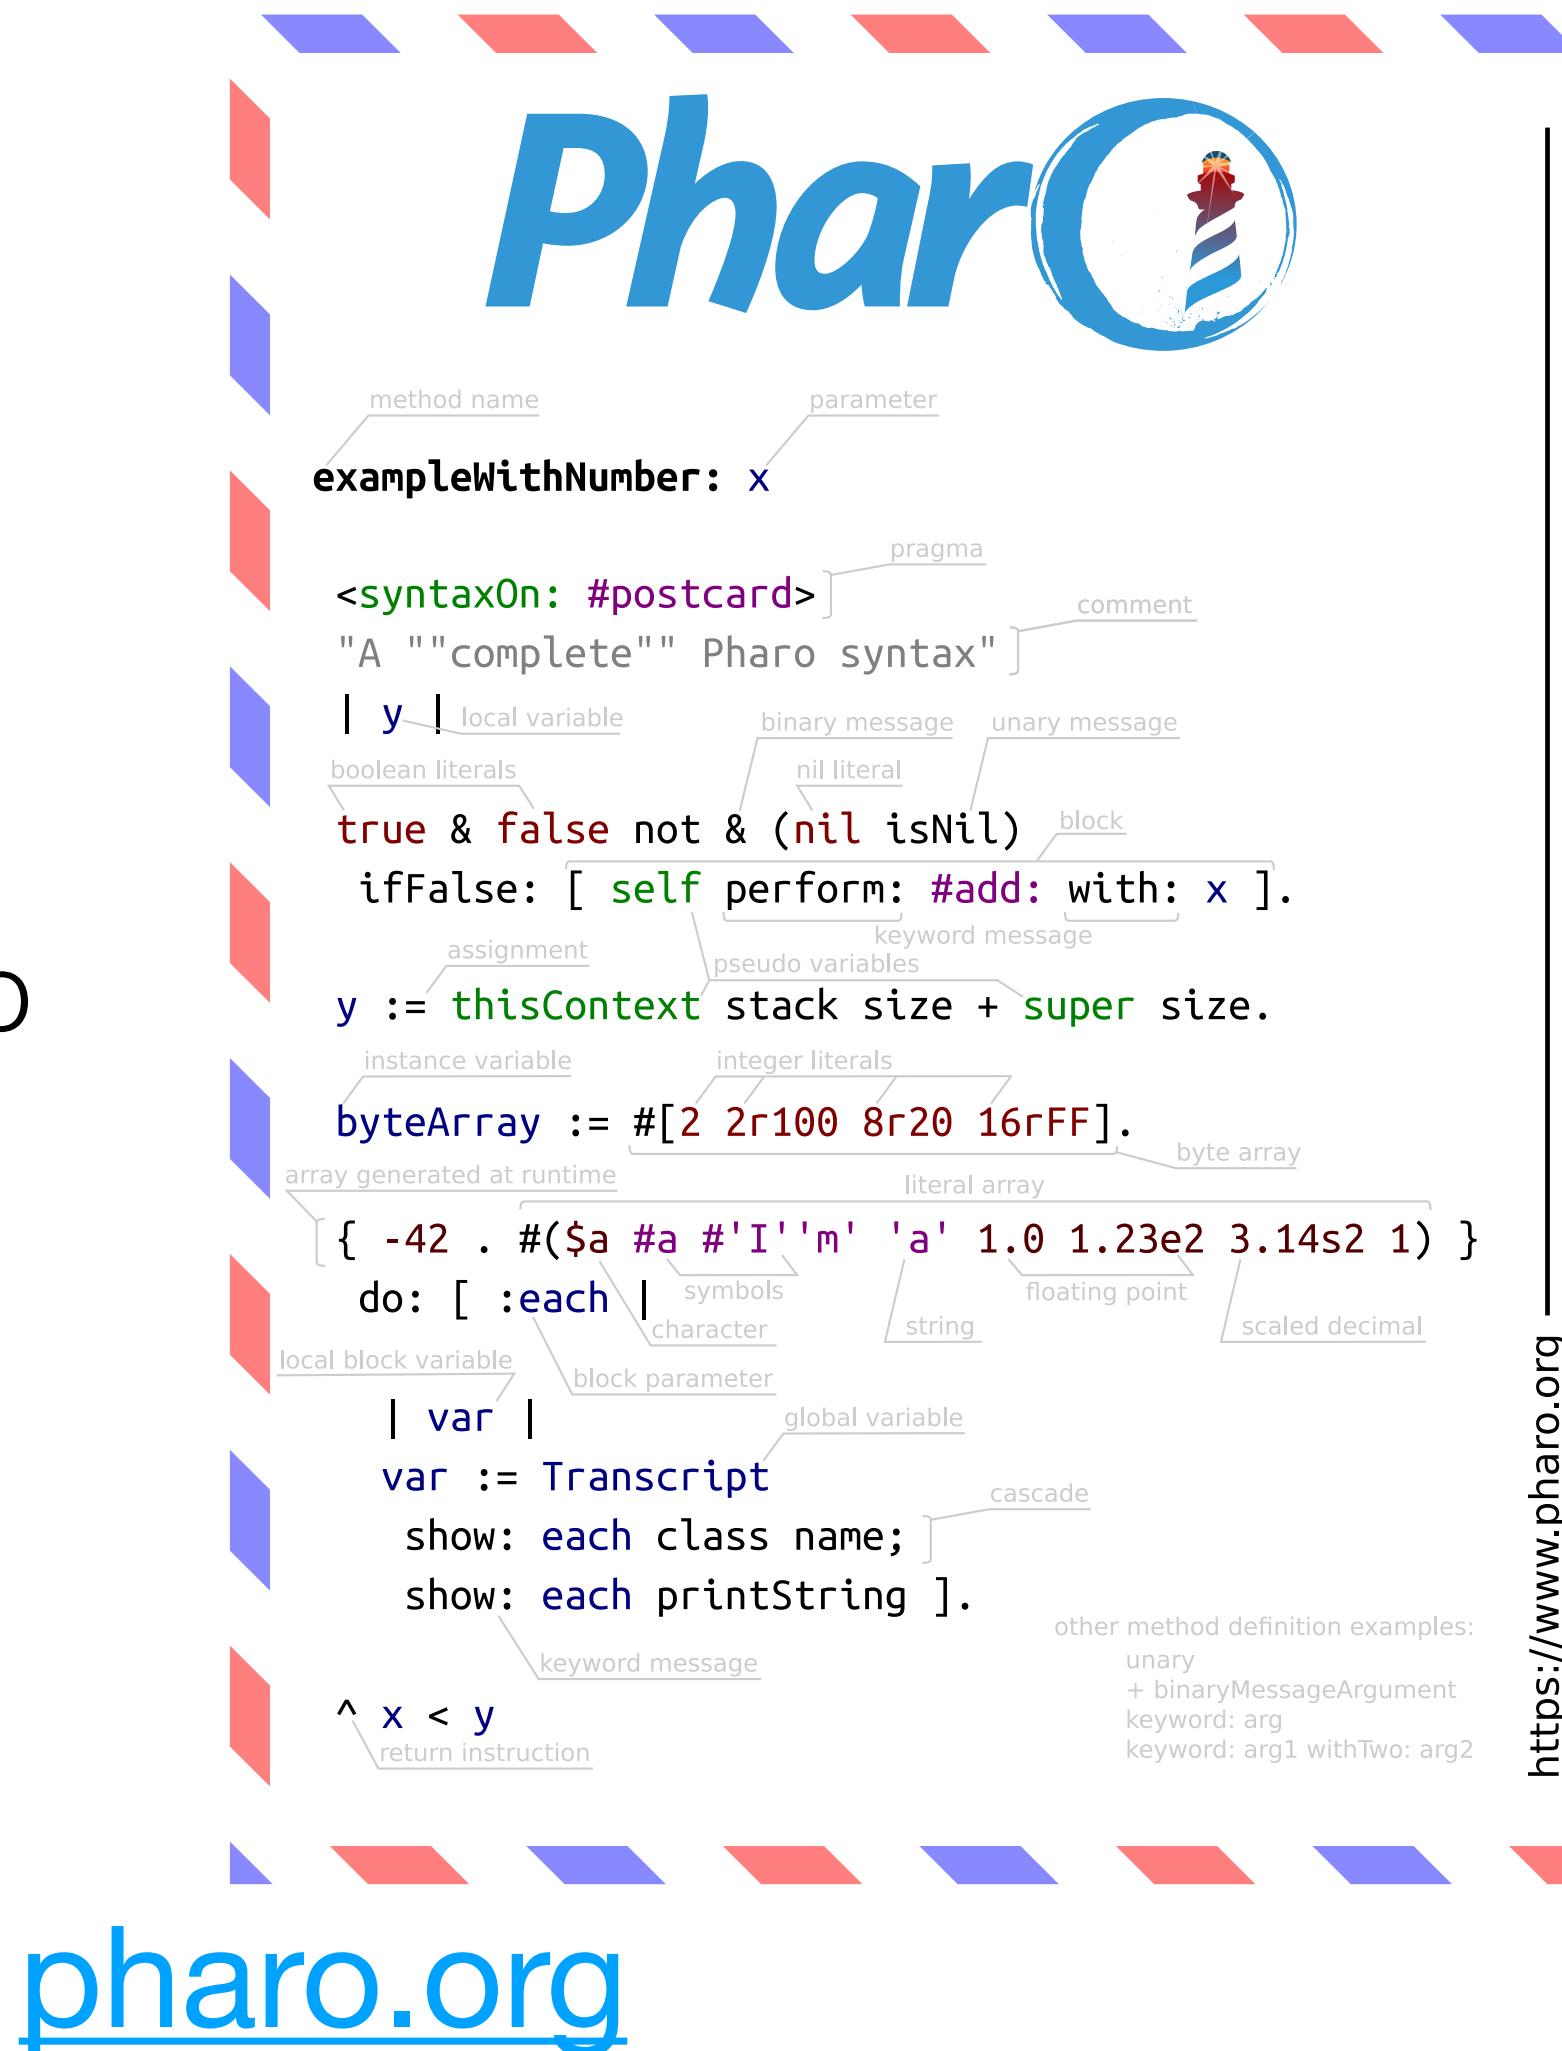
- **Bachelor's:** Software Engineer, UCB, Bolivia
- **Master's:** Software Engineering, UL, France
- **PhD:** Started in 2023 in memory profiling
- **Interests:** Music (progressive rock, Charly), (real) languages, sports, Pharo



Pharo: a programming language



- Advanced run time reflection
- Pure object-oriented approach
- Software as objects
- (Super) simple syntax
- Pharo, and its VM, are written in Pharo
- MIT open-source license
- Live inspectable objects
- Easy stack manipulation

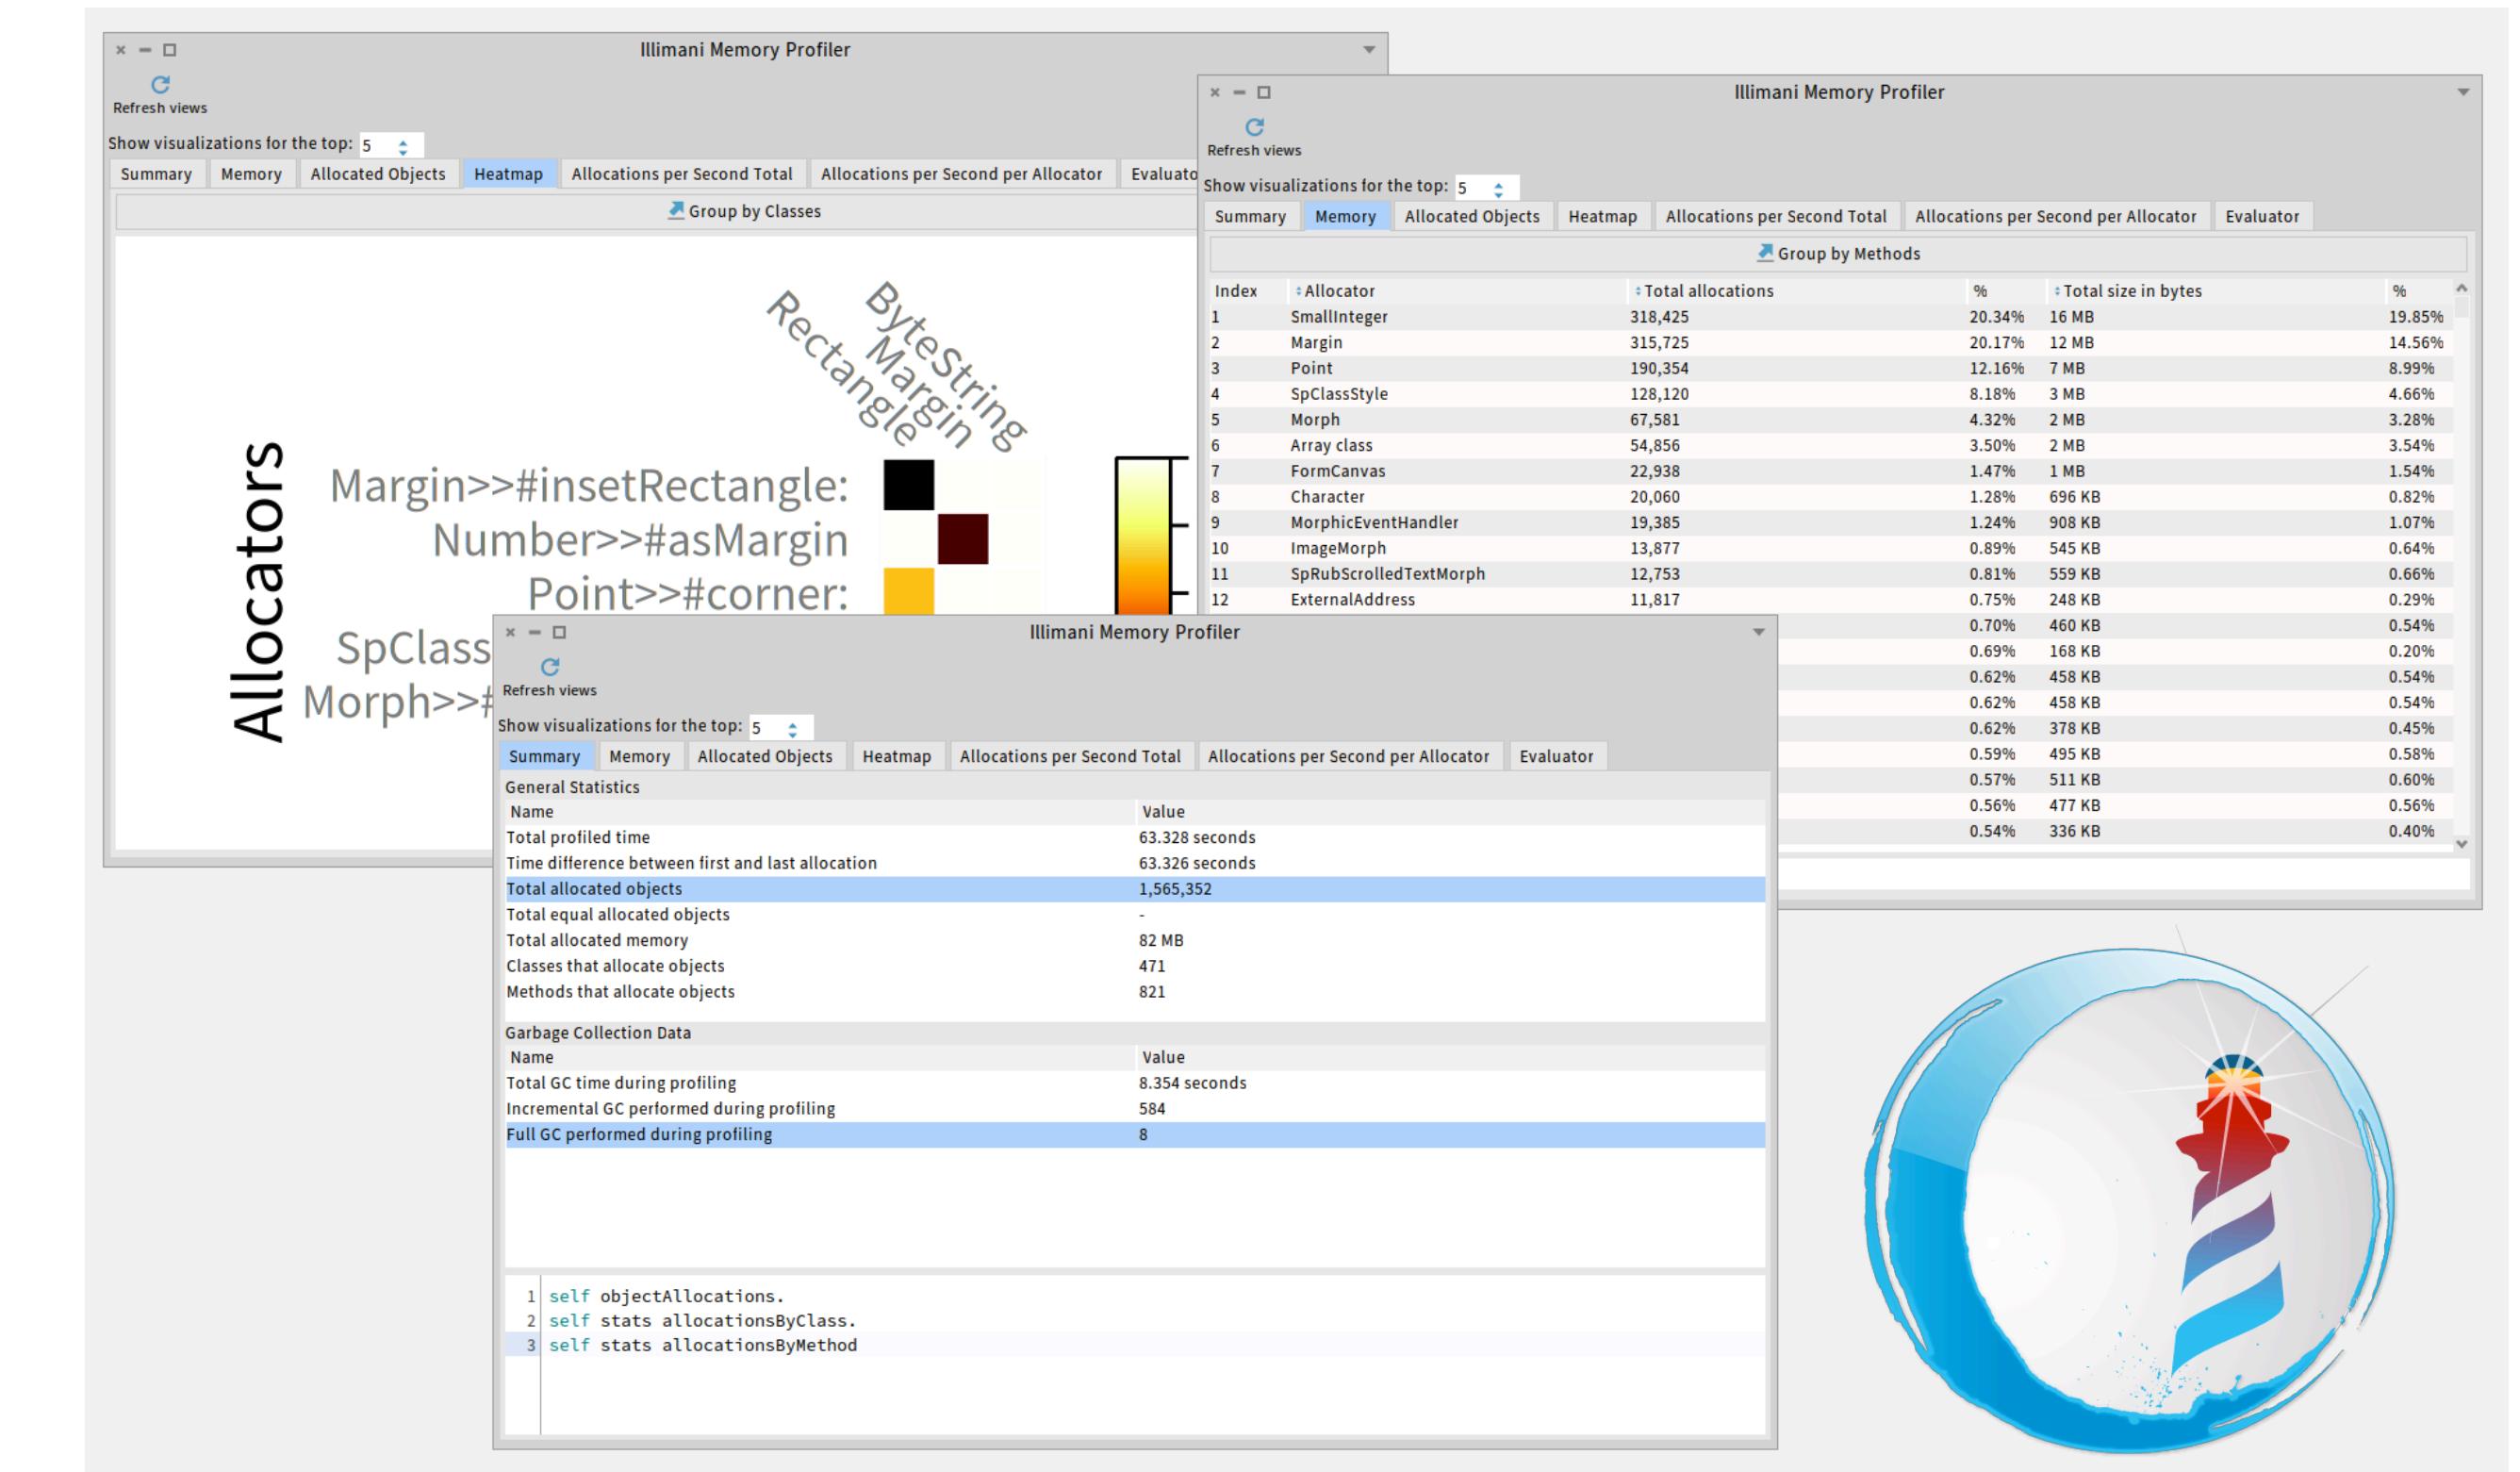


PLACE
STAMP
HERE

<https://www.pharo.org>

Illimani: a Pharo memory profiler

- Open-source MIT license
- Detects object allocation sites
- Tracks object lifetimes
- Allocation matrix
- Unmodified VM
- Density chart
- Memory consumption tables
- Rich object-oriented model



github.com/jordanmontt/illimani-memory-profiler

Object allocation sites

We define an object allocation site as the textual location in the source code where the object was created [1]

AthensTextScanner >> initialize

Allocation site

lines := **OrderedCollection** new

...



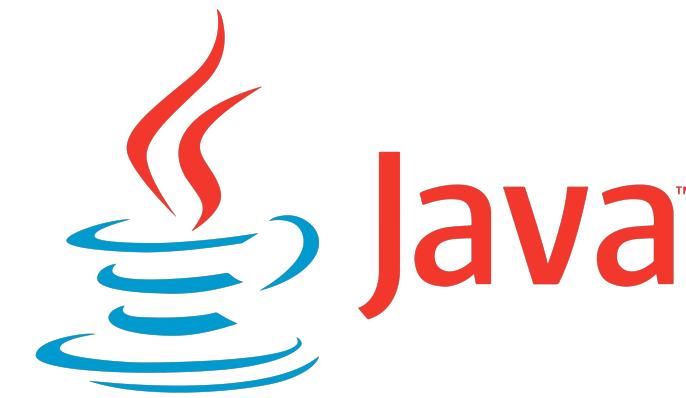
[1]

Memento Mori: Dynamic Allocation-Site-Based Optimizations

Daniel Clifford Hannes Payer Michael Stanton Ben L. Titzer

Capture object allocation sites

In Pharo, almost all computations are done by sending a message. This is also true when allocating objects.



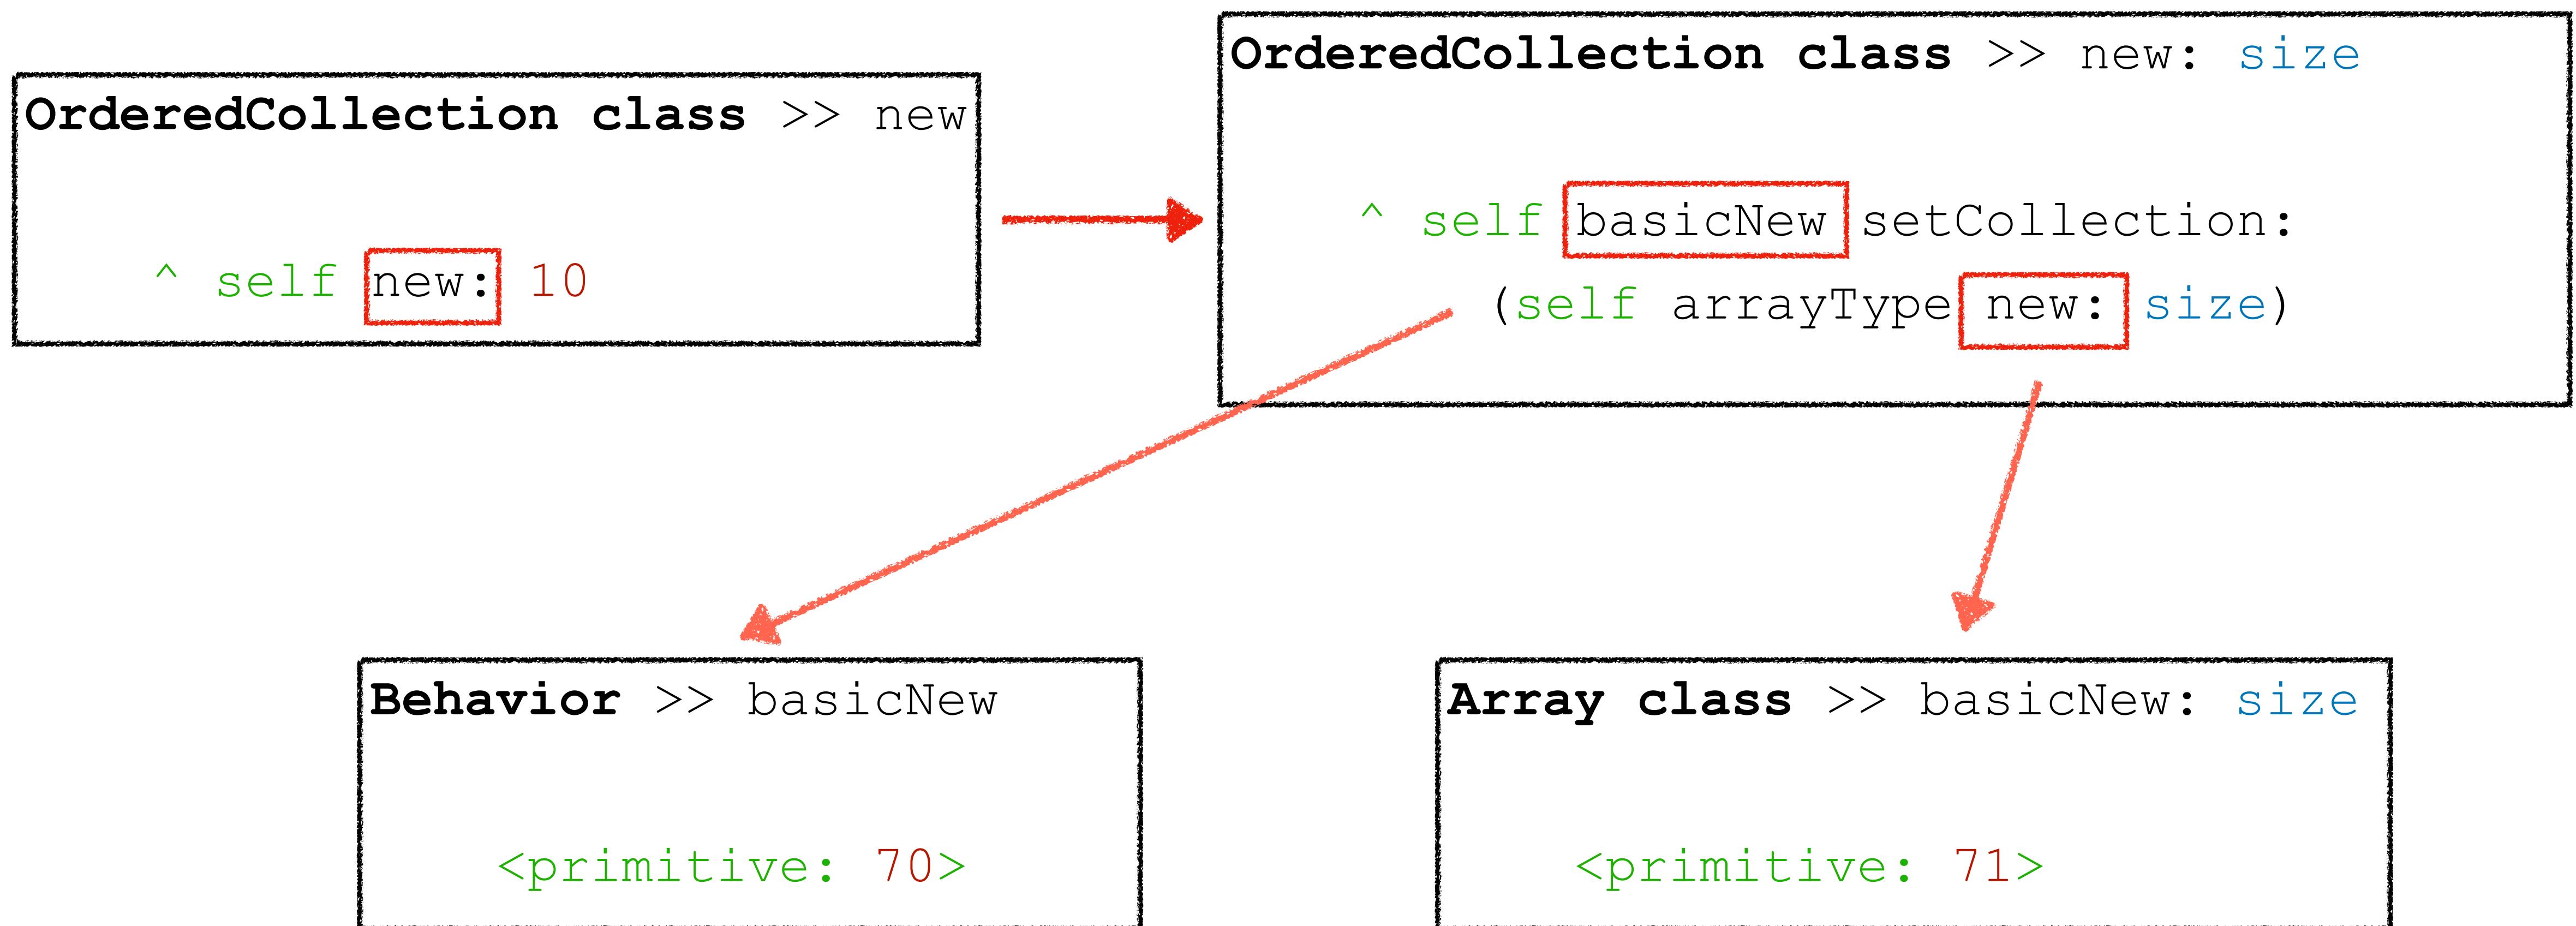
```
collection := OrderedCollection new.
```

```
list = new ArrayList();
```

Allocating an object

```
OrderedCollection new
```

Allocating an object



Capture object allocation sites

```
AthensTextScanner >> initialize
```

```
lines := OrderedCollection new
```

```
...
```

```
OrderedCollection class >> new: size
```

```
^ self basicNew setCollection:  
(self arrayType new: size)
```

```
Behavior >> basicNew
```

```
<primitive: 70>
```

Instrumentation

Allocation site

```
Array class >> basicNew: size
```

```
<primitive: 71>
```

Allocating an object

```
AthensTextScanner >> initialize
```

```
lines := OrderedCollection new
```

```
...
```

```
OrderedCollection class >> new: size
```

```
^ self basicNew setCollection:  
(self arrayType new: size)
```

```
Behavior >> basicNew
```

```
<primitive: 70>
```

```
Array class >> basicNew: size
```

```
<primitive: 71>
```

Instrumentation

Allocation site

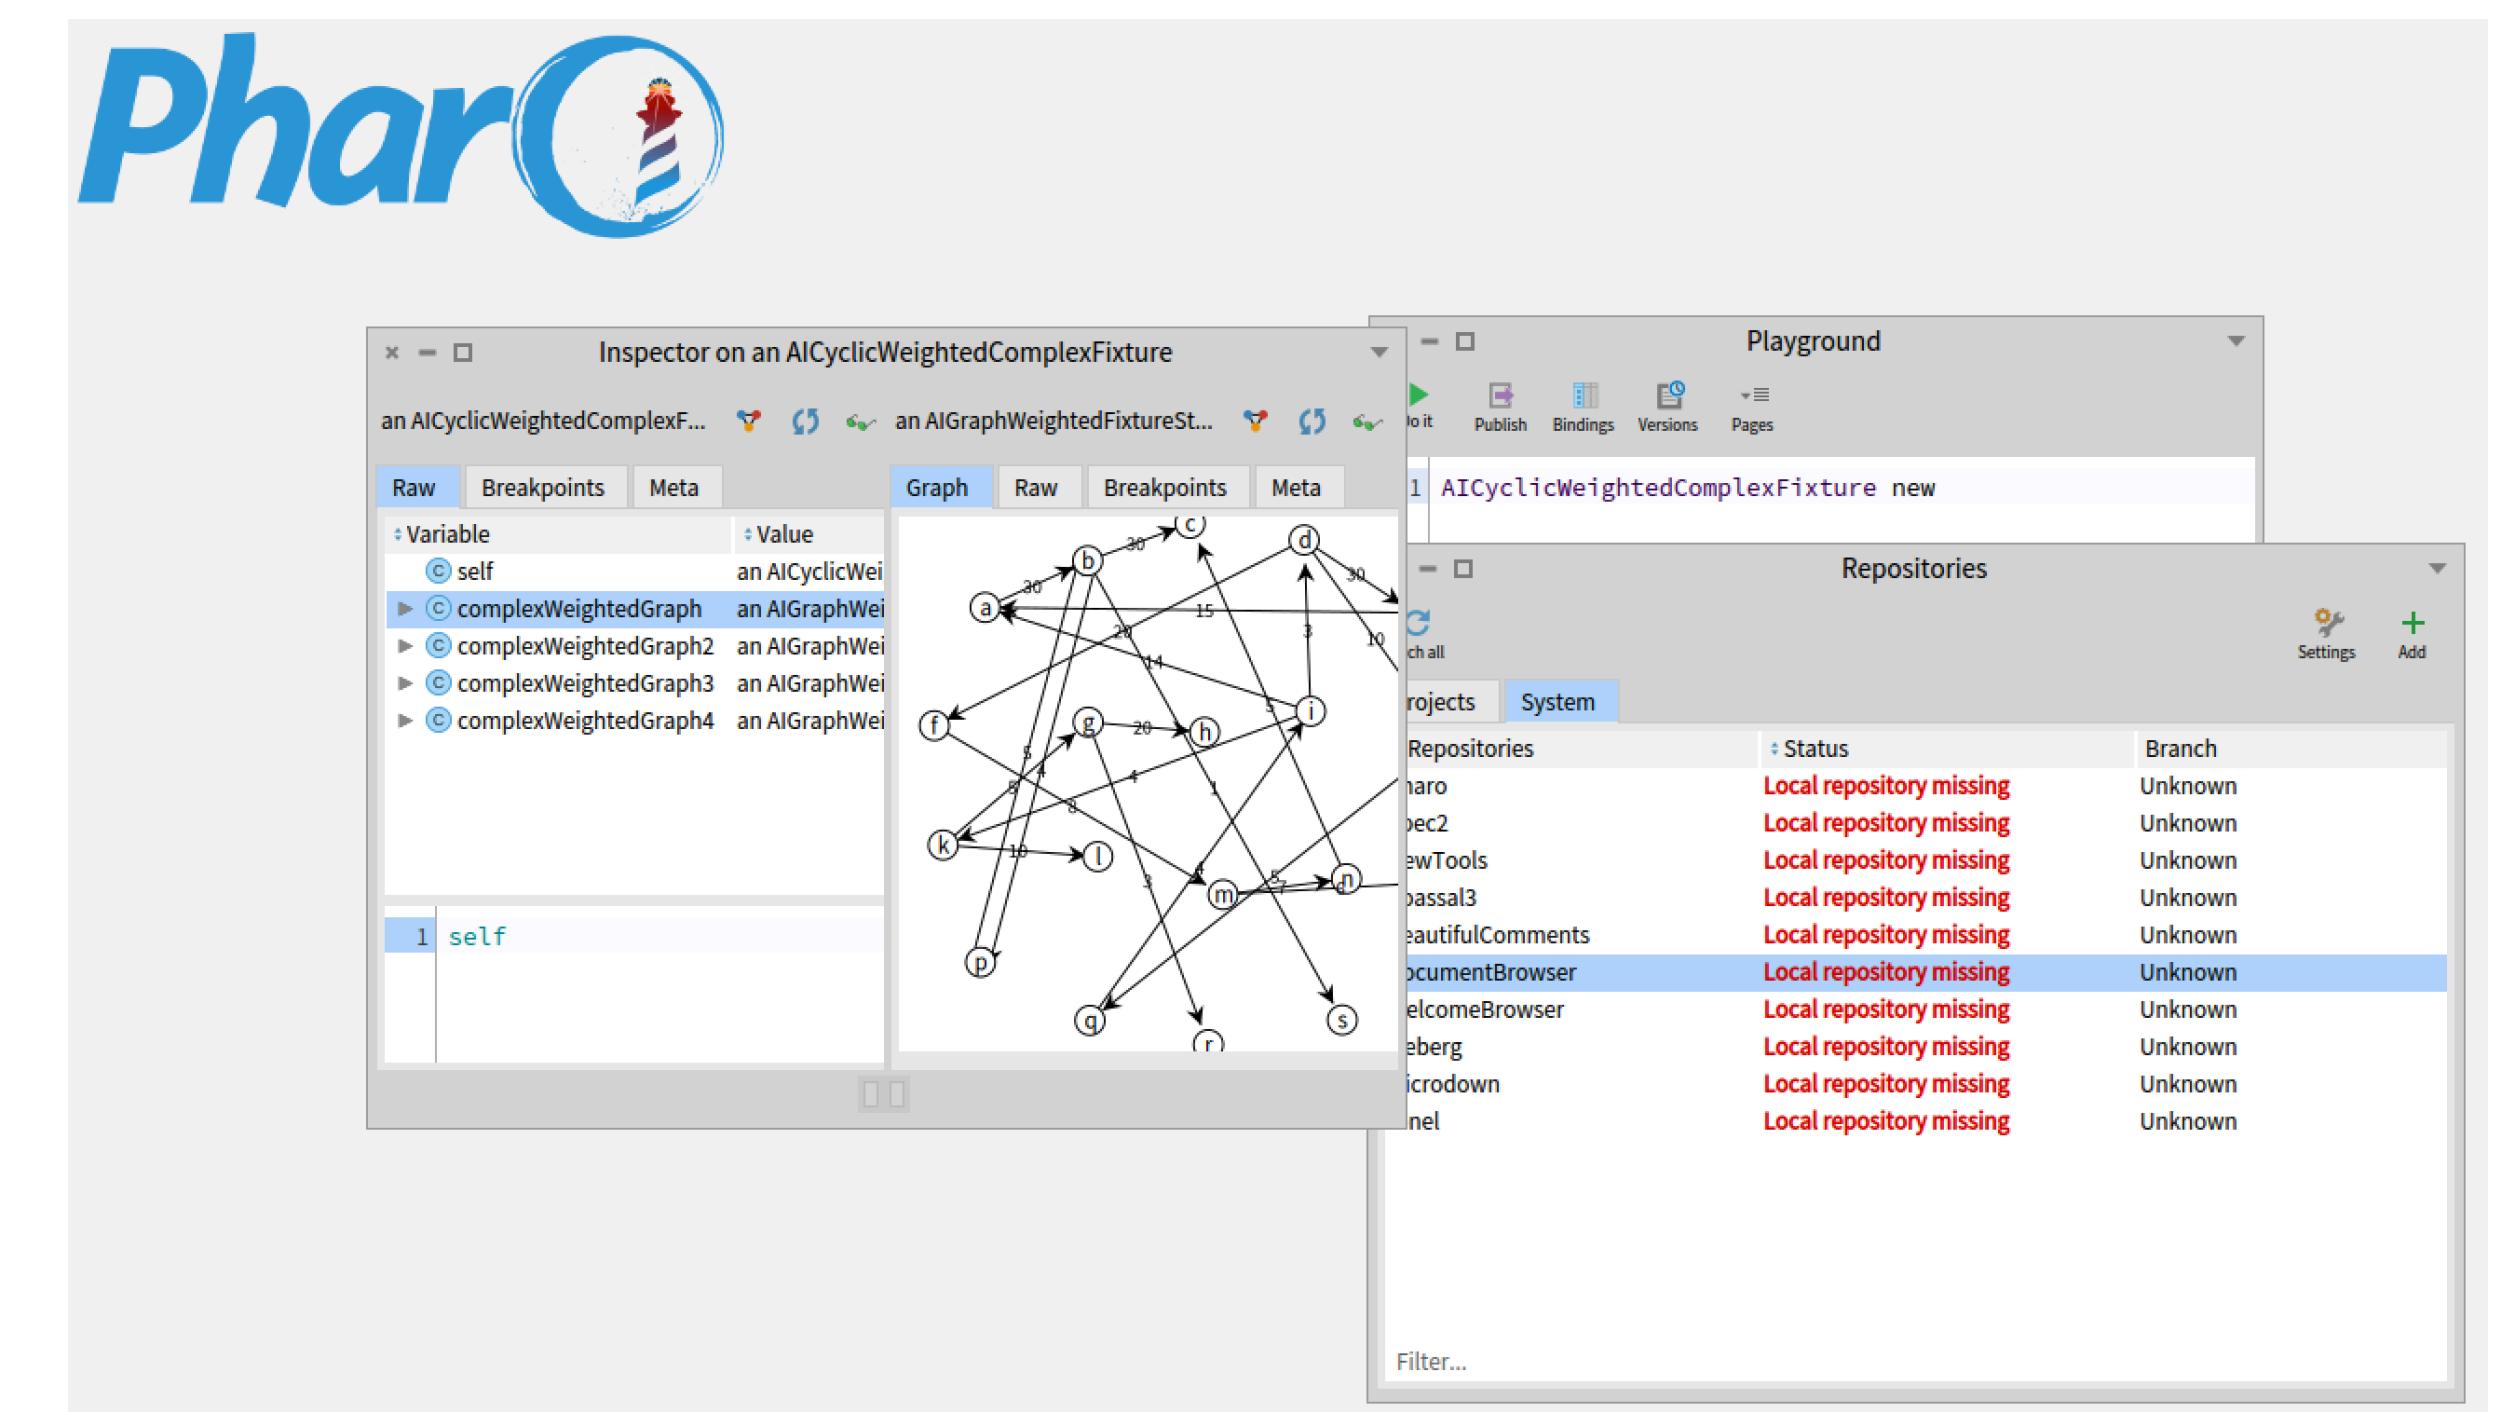
Instrumenting object allocation sites in Pharo

We've instrumented 7 methods that allocate objects to capture when they will be called and then filter the execution stack.

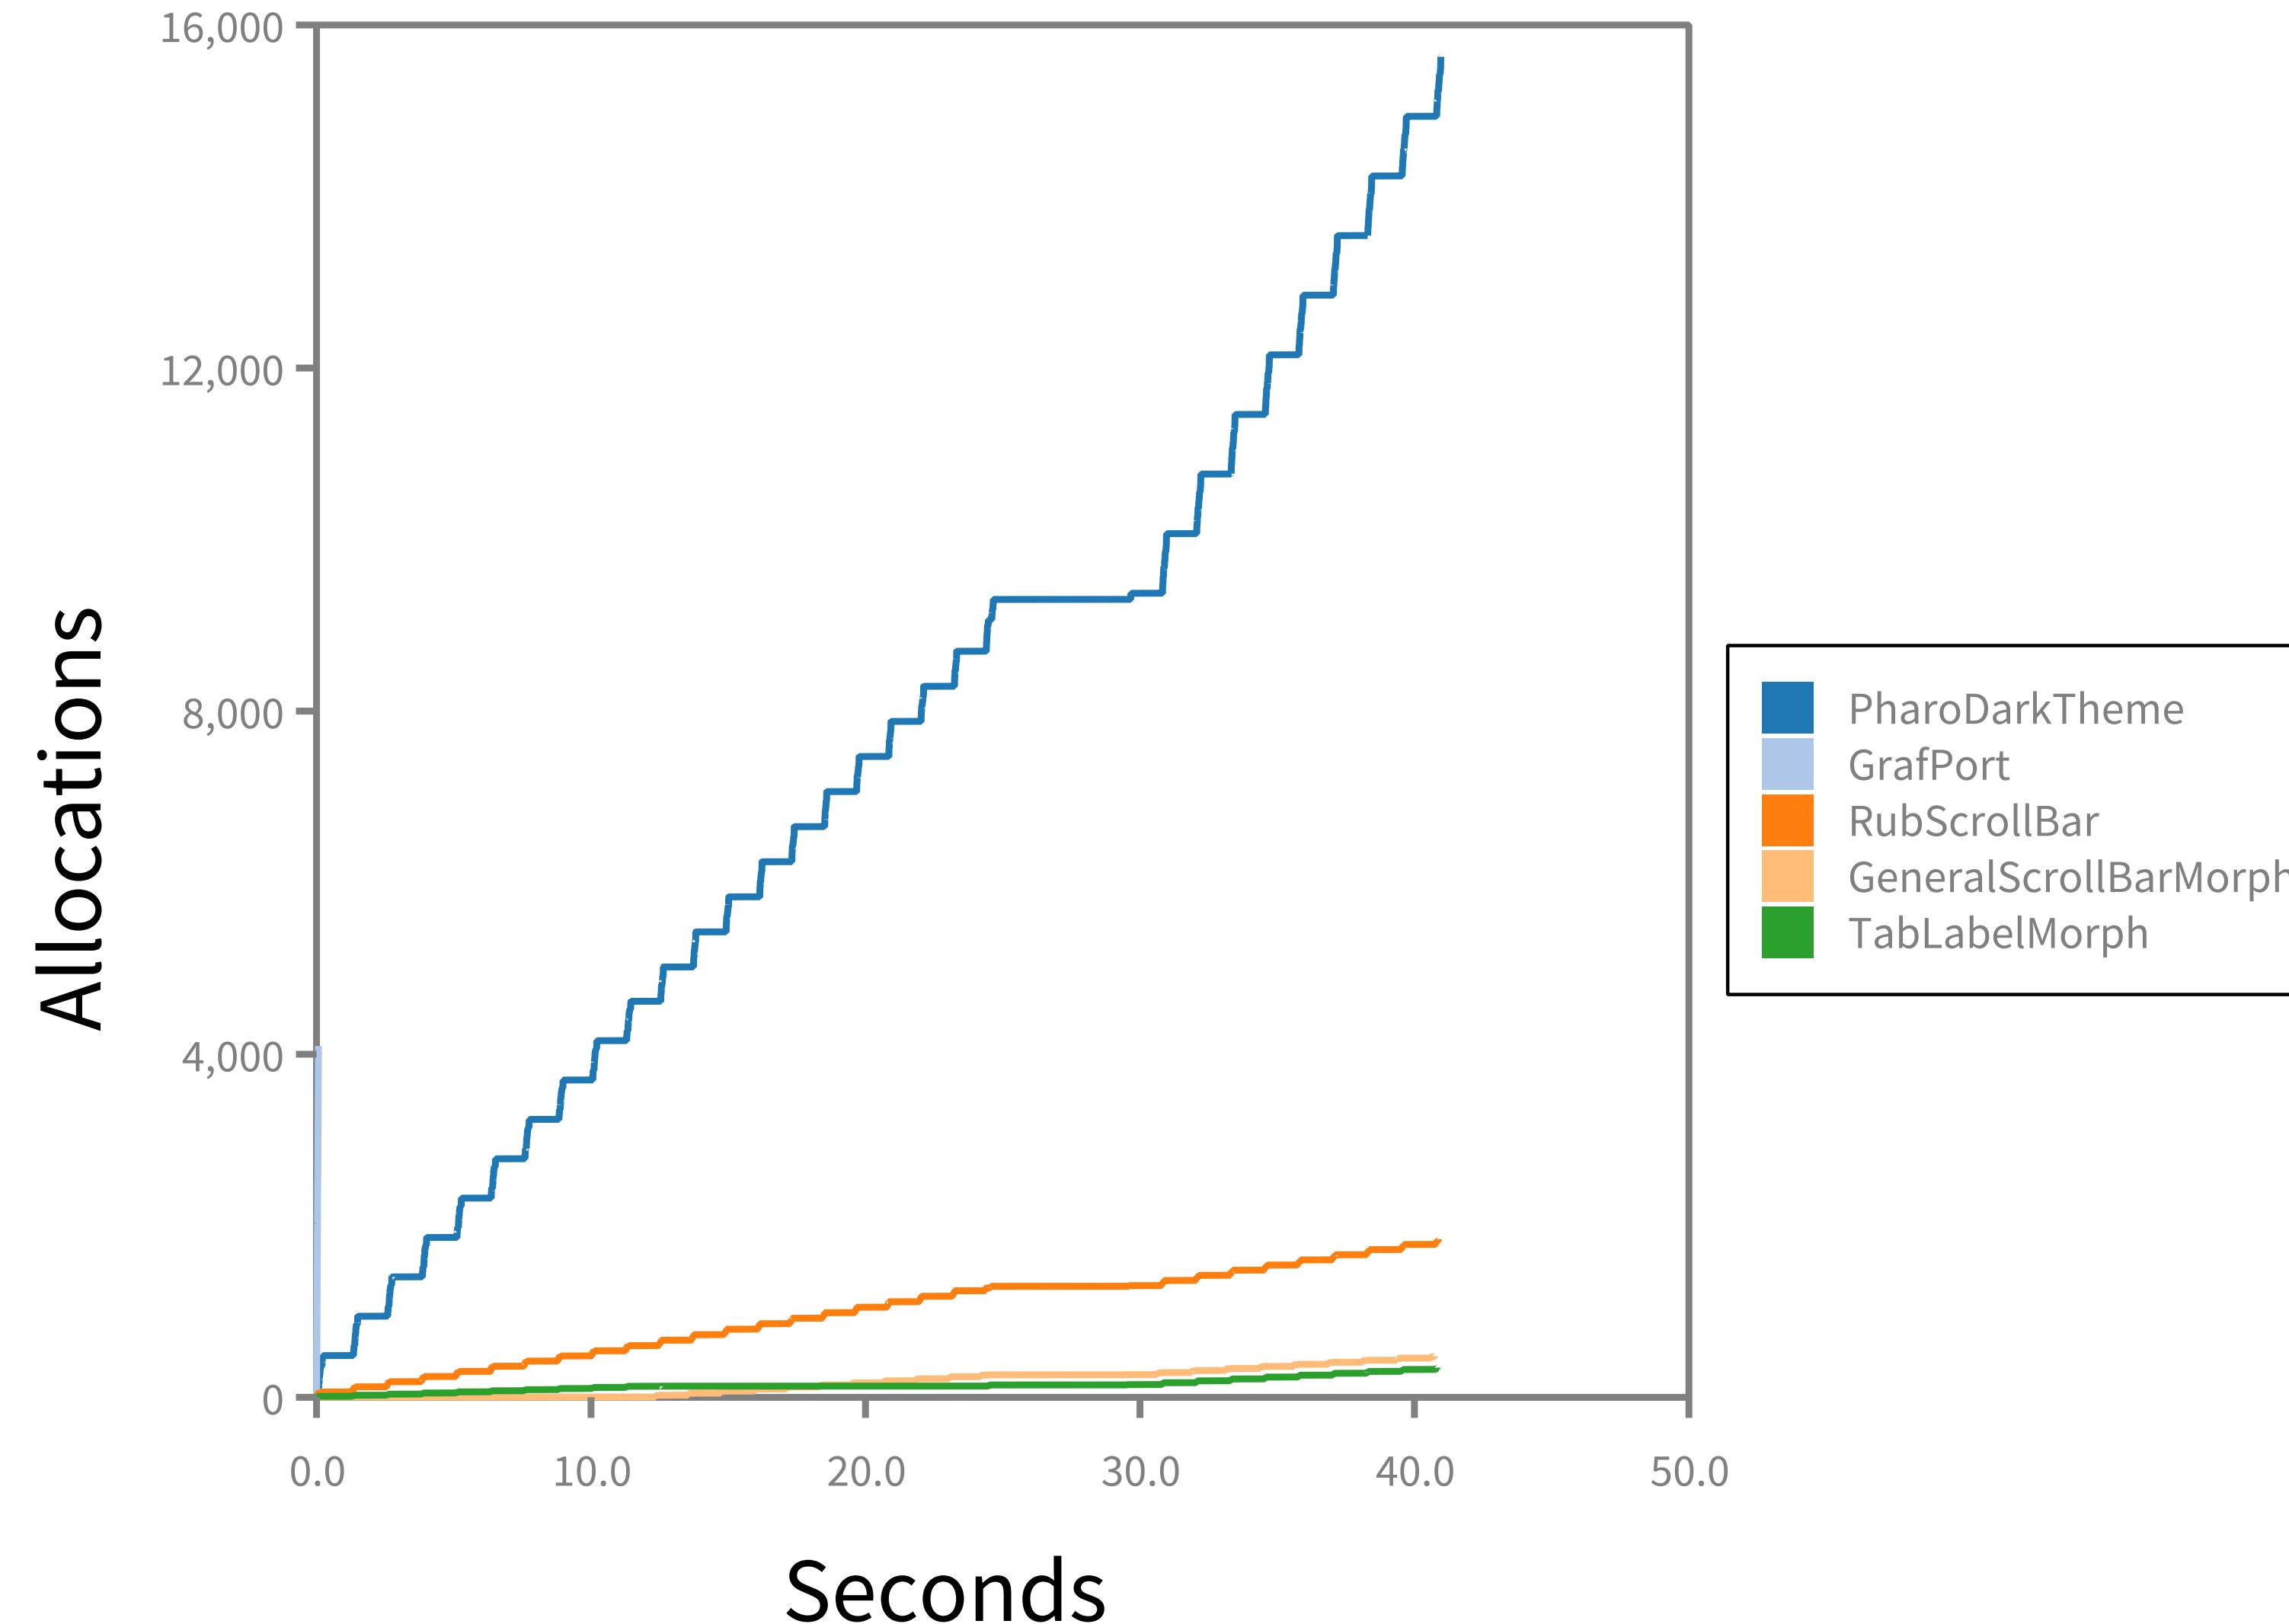
- Behavior >> `basicNew`
- Behavior >> `basicNew:`
- Behavior >> `basicPinned`
- Behavior >> `basicPinned:`
- Behavior >> `basicOld`
- Behavior >> `basicOld:`
- Array class >> `new:`

Case study 1: Morphic

- We profiled the allocations of type **Color**
- We opened 30 Pharo tools (Inspector, Playground and Iceberg) and we let them render for 100 rendering cycles



Allocations over time

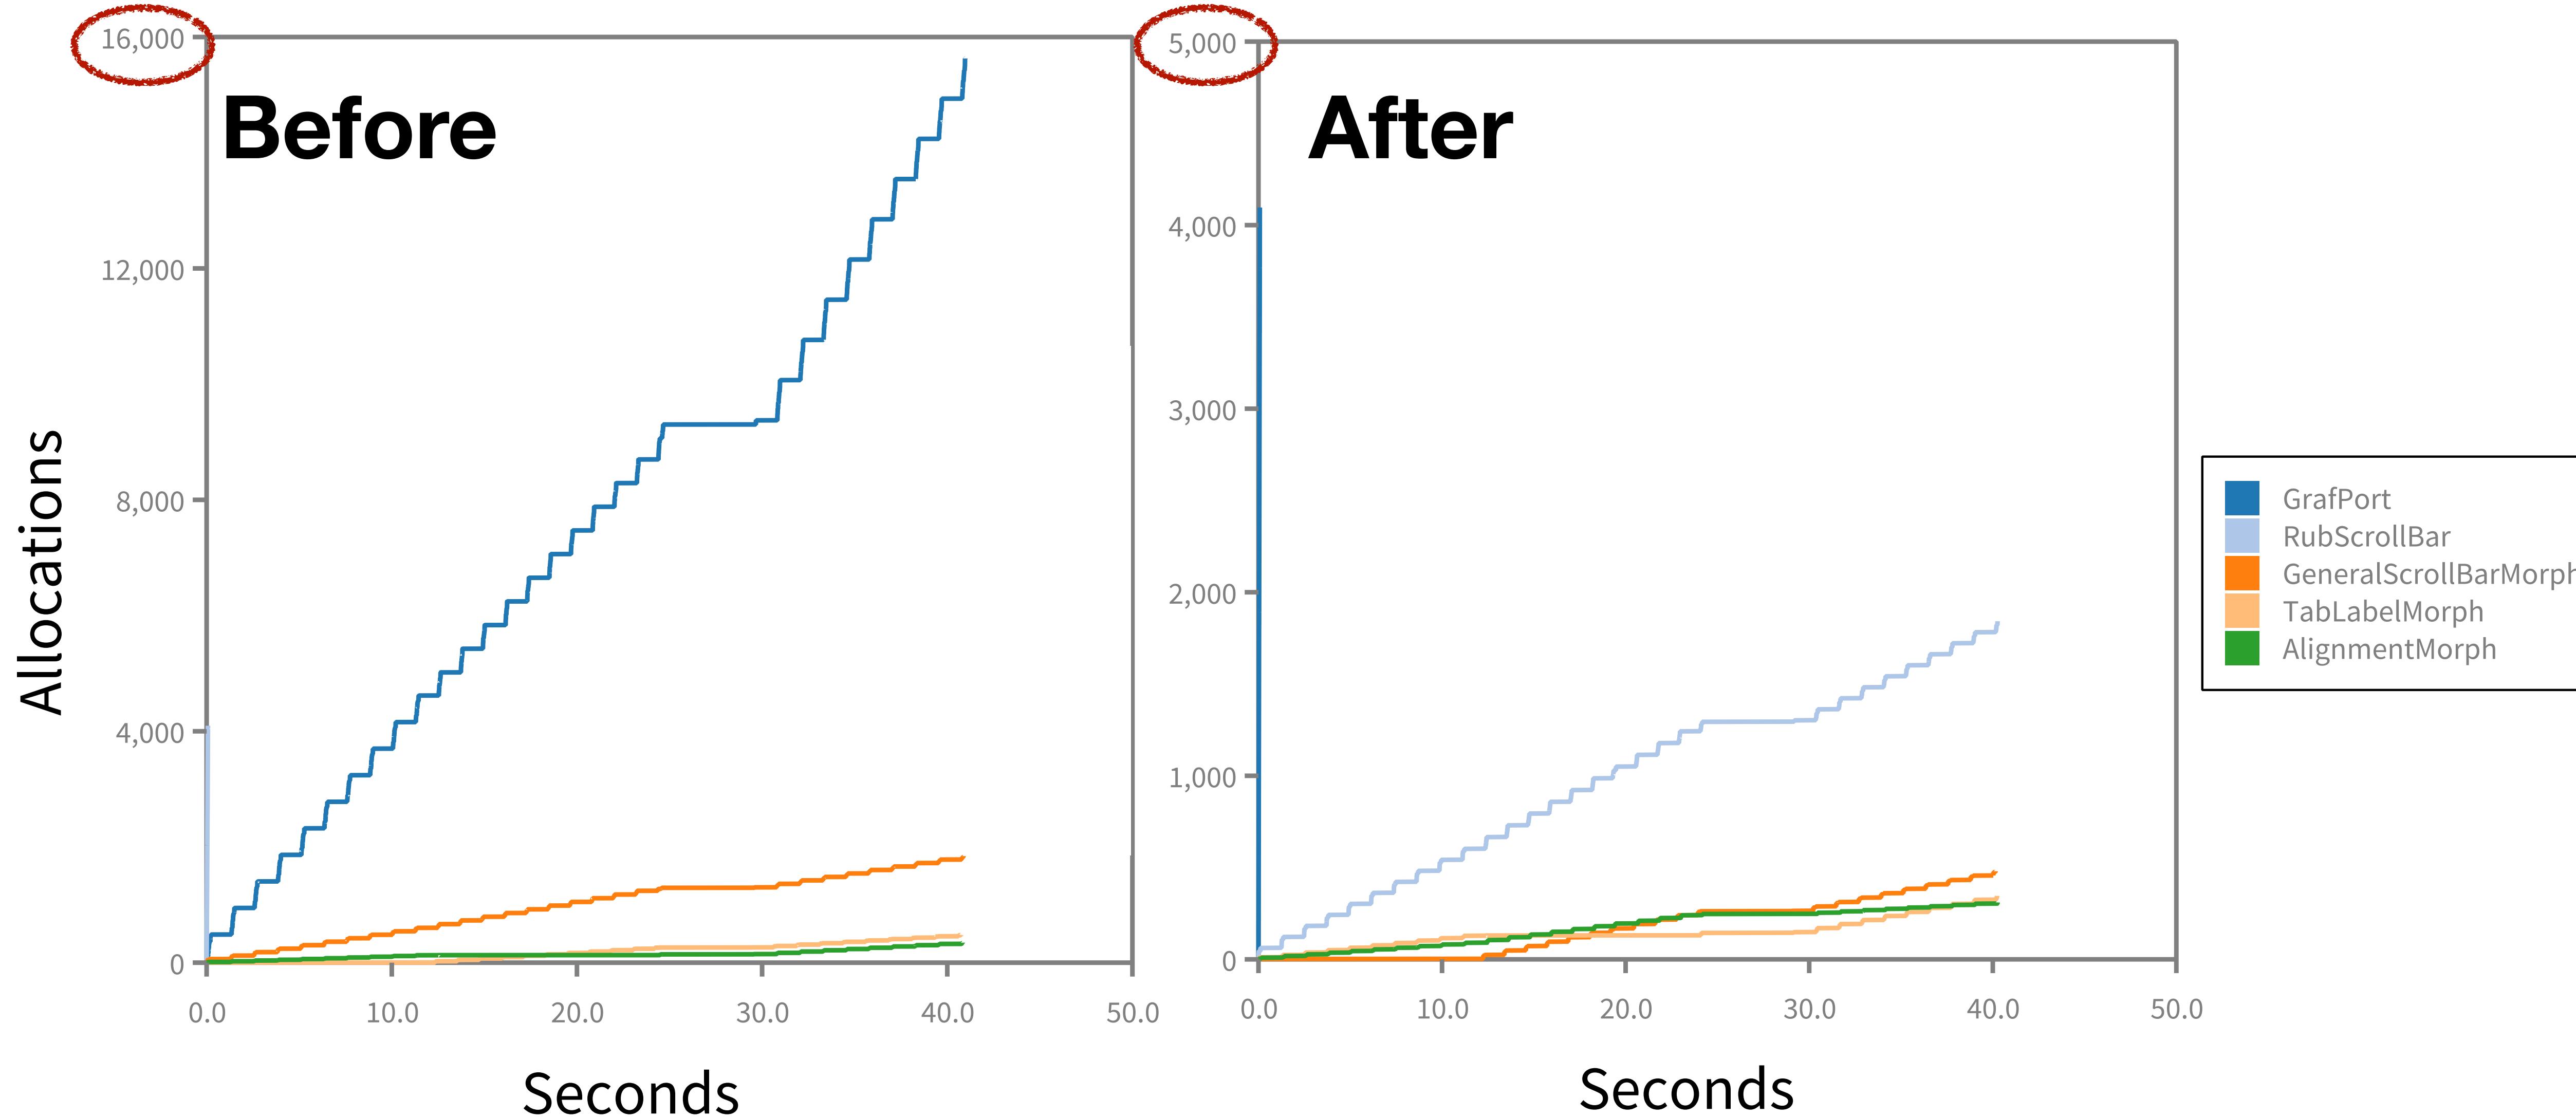


Morphic Analysis

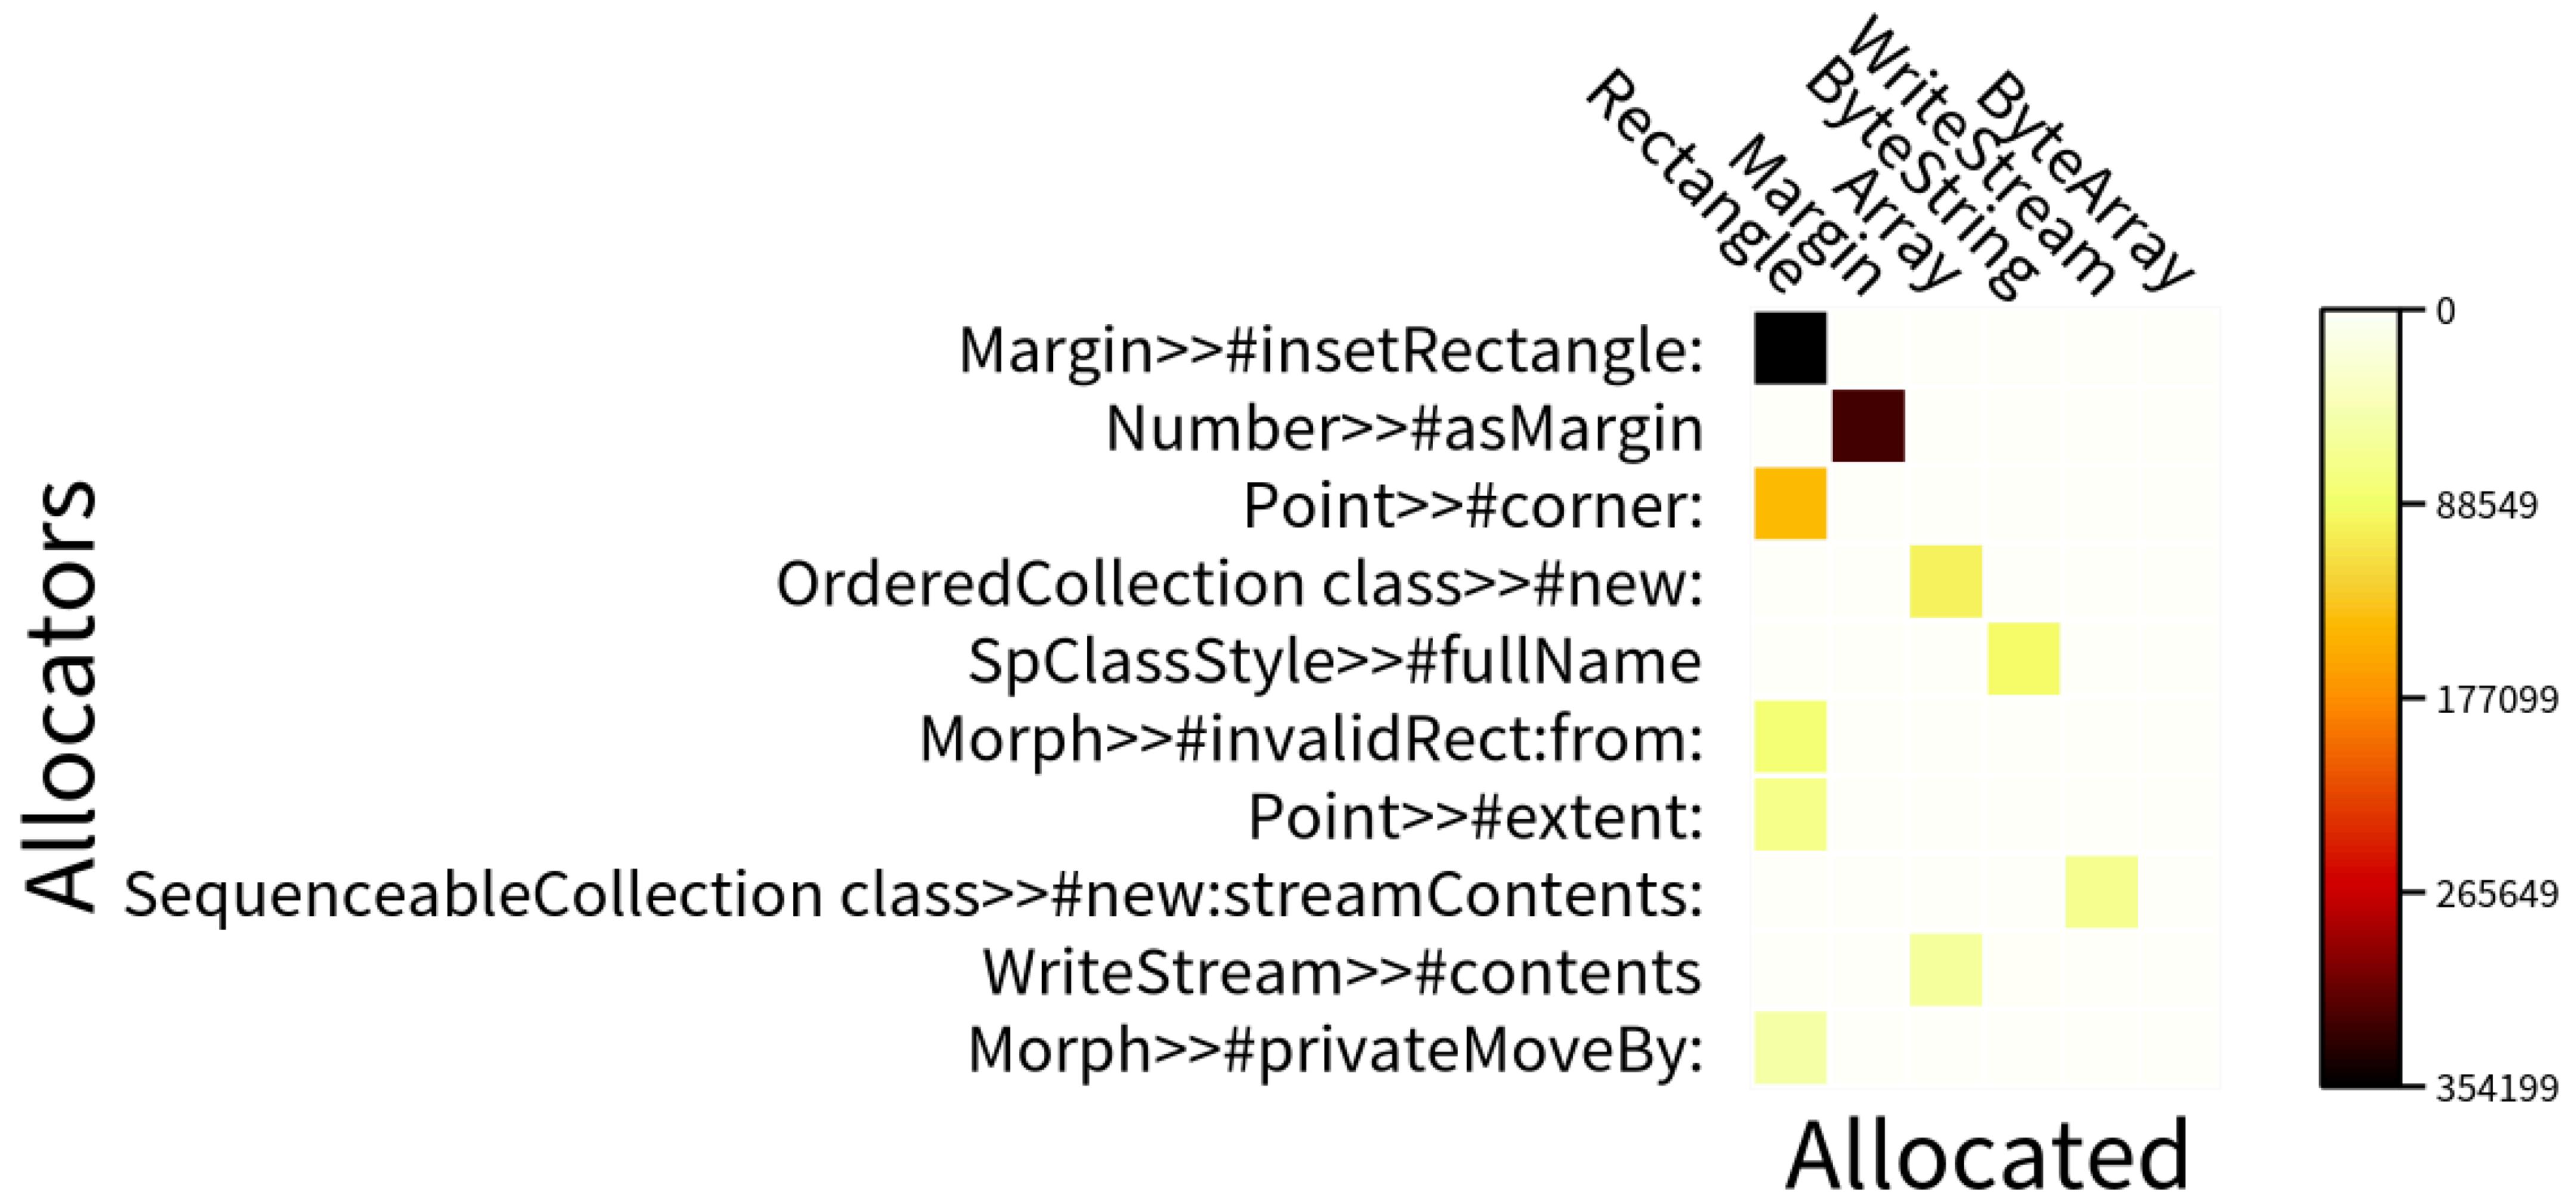
| Allocator class | Allocated colors | % |
|-----------------------|------------------|-----|
| PharoDarkTheme | 15,629 | 66% |
| GrafPort | 4,096 | 17% |
| RubScrollBar | 1,842 | 8% |
| GeneralScrollBarMorph | 480 | 2% |
| TabLabelMorph | 346 | 1% |
| Rest of the classes | 1293 | 2% |

We have identified an object allocation site in the class `PharoDarkTheme` that allocated 66% of all the allocated colors with 99,9% redundant allocations.

Morphic after the fix

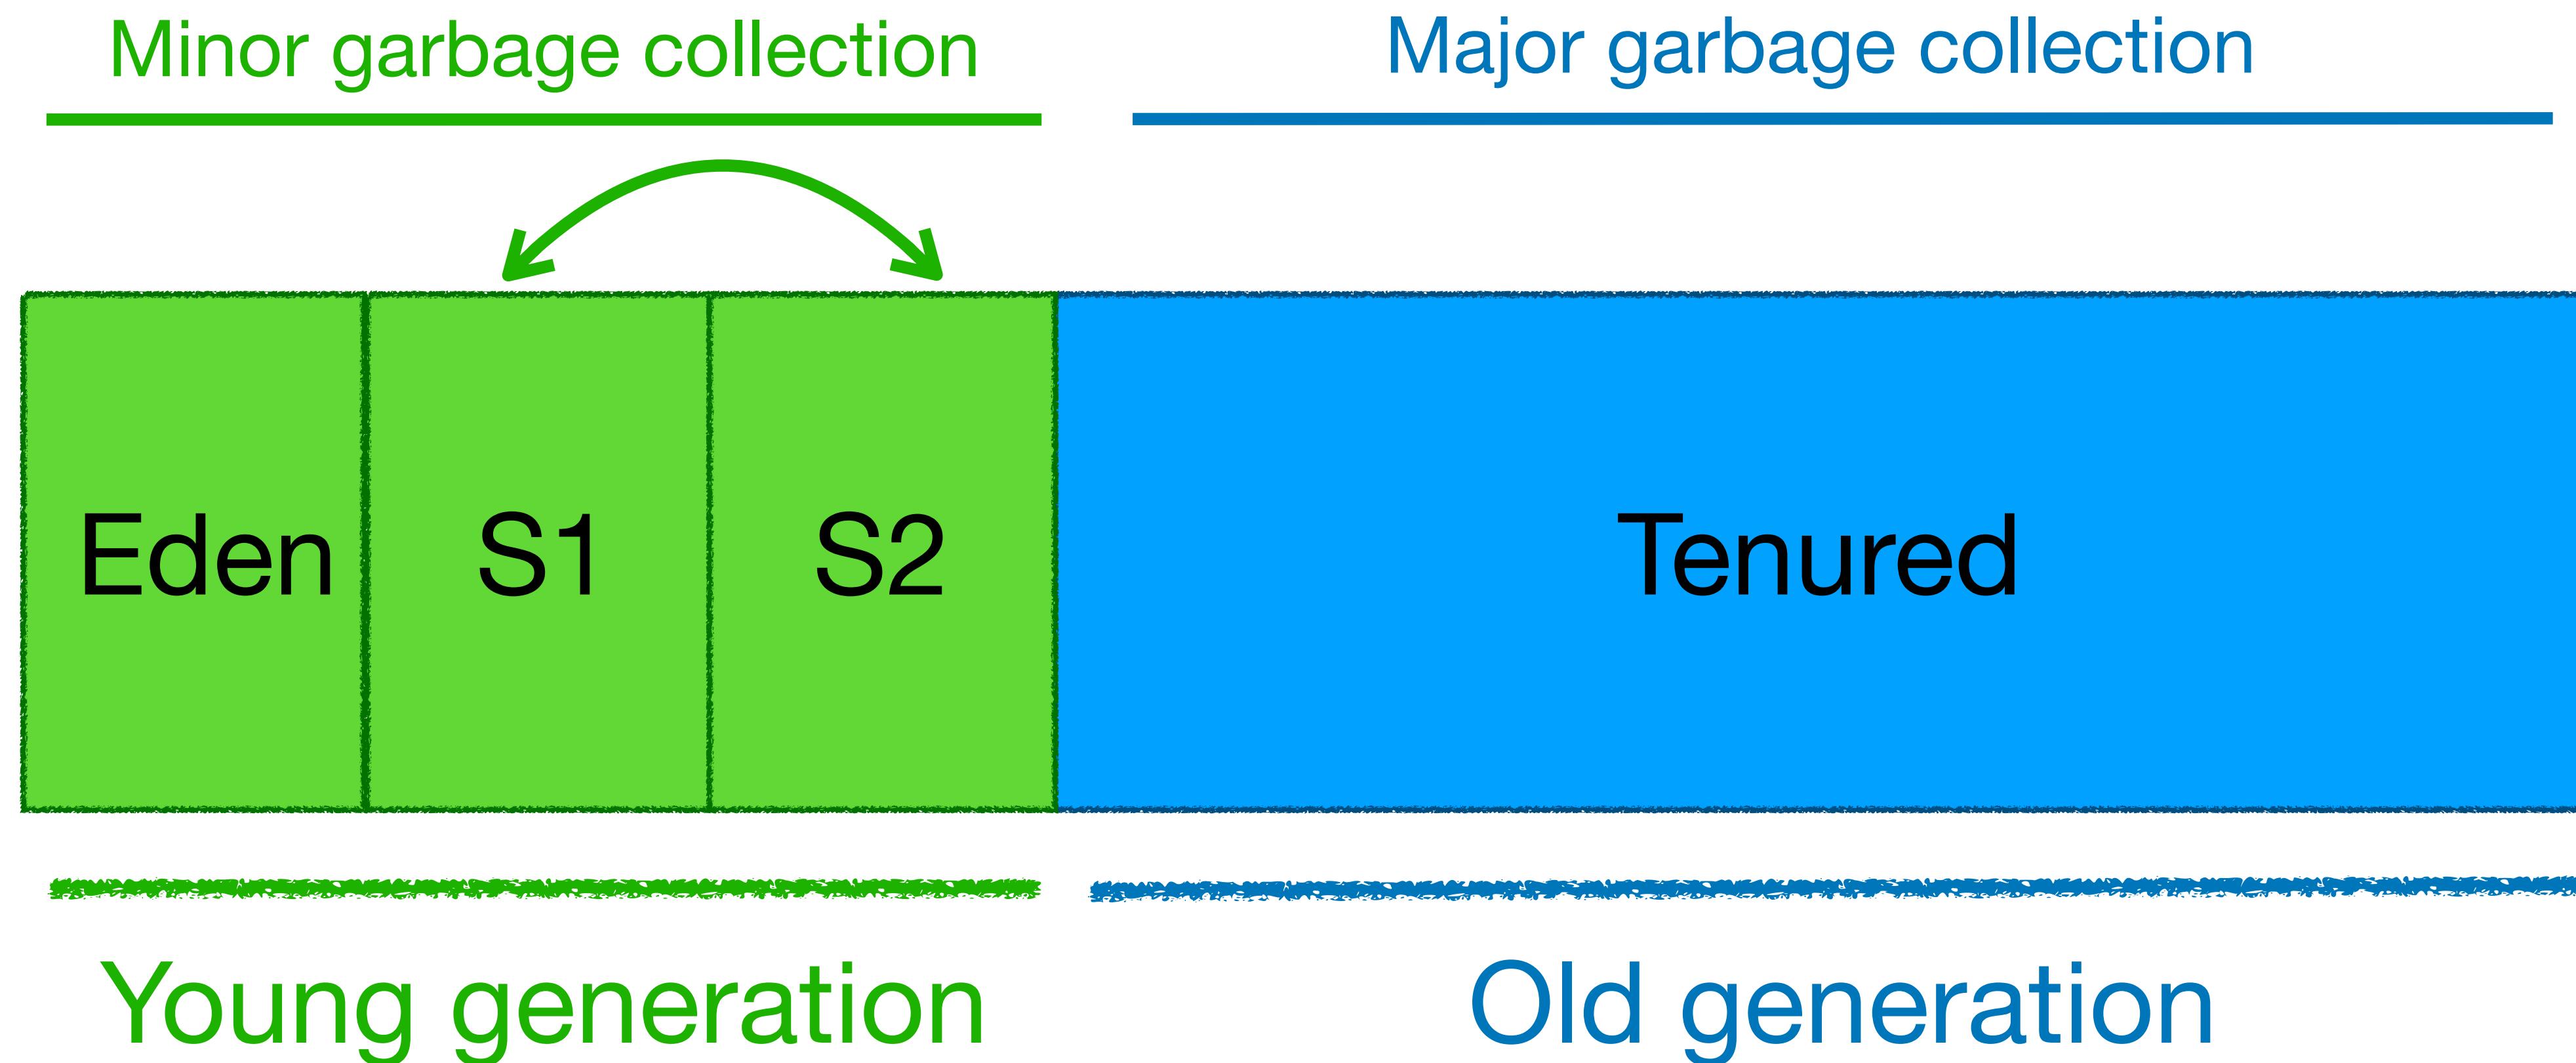


Detecting other allocation sites



Case study 2: Object lifetimes

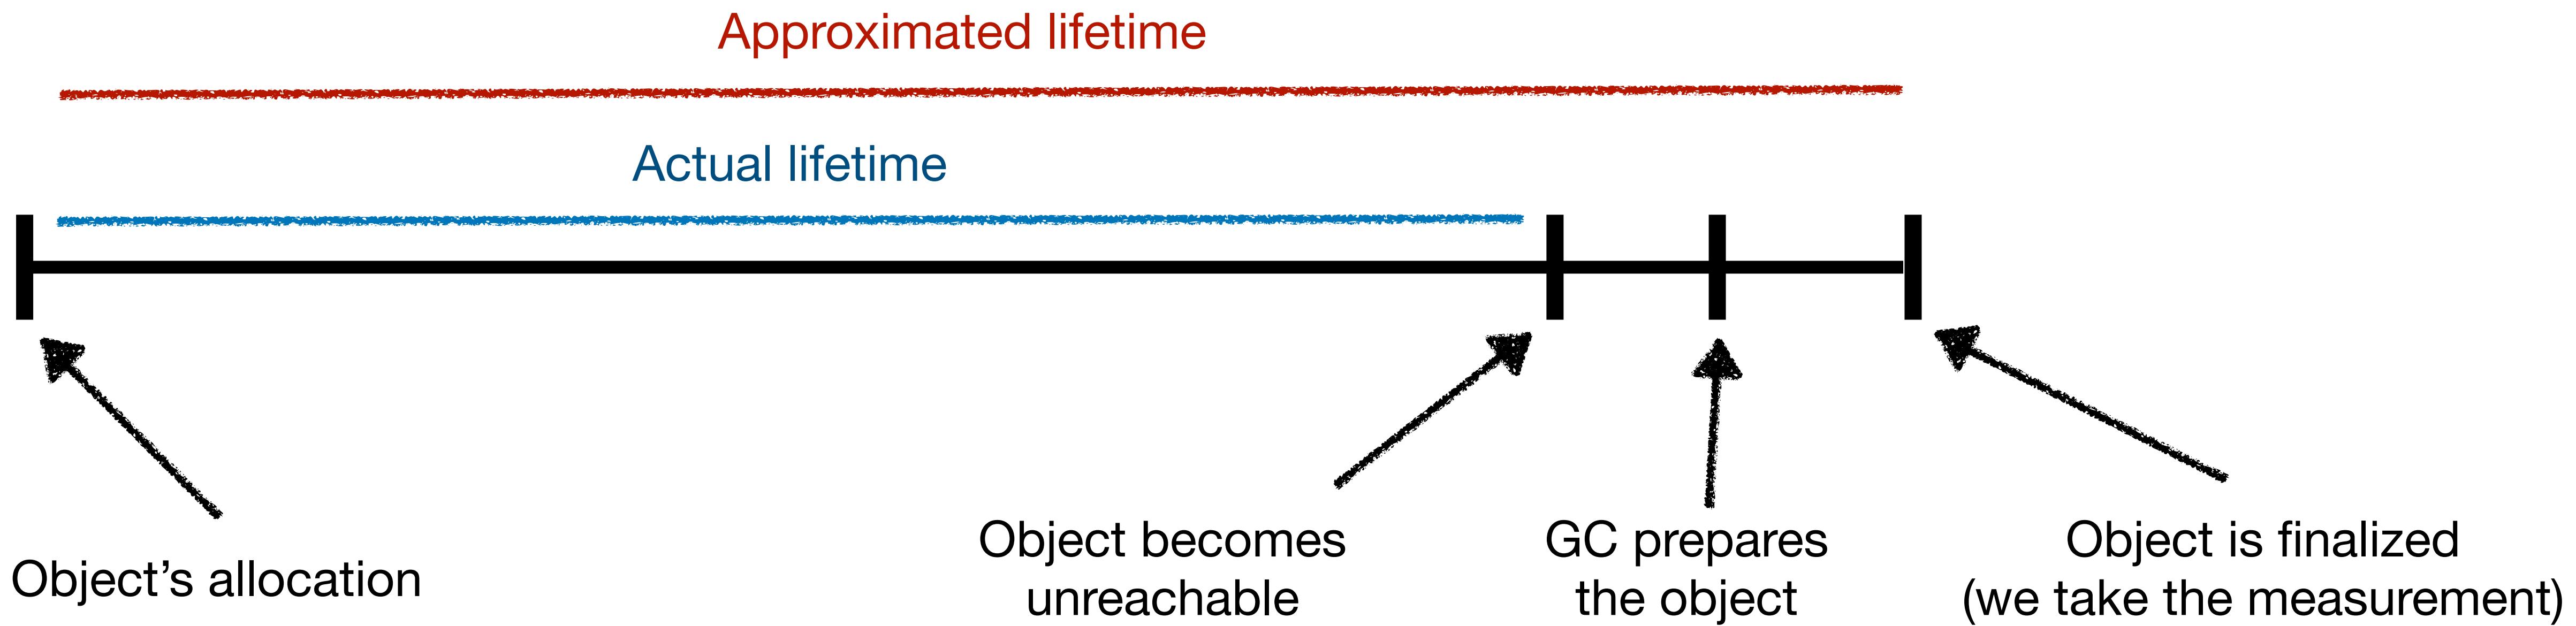
Pharo's garbage collector



An object's (approximated) lifetime

$$\text{object'sLifetime} = \text{finalizationTime} - \text{allocationTime}$$

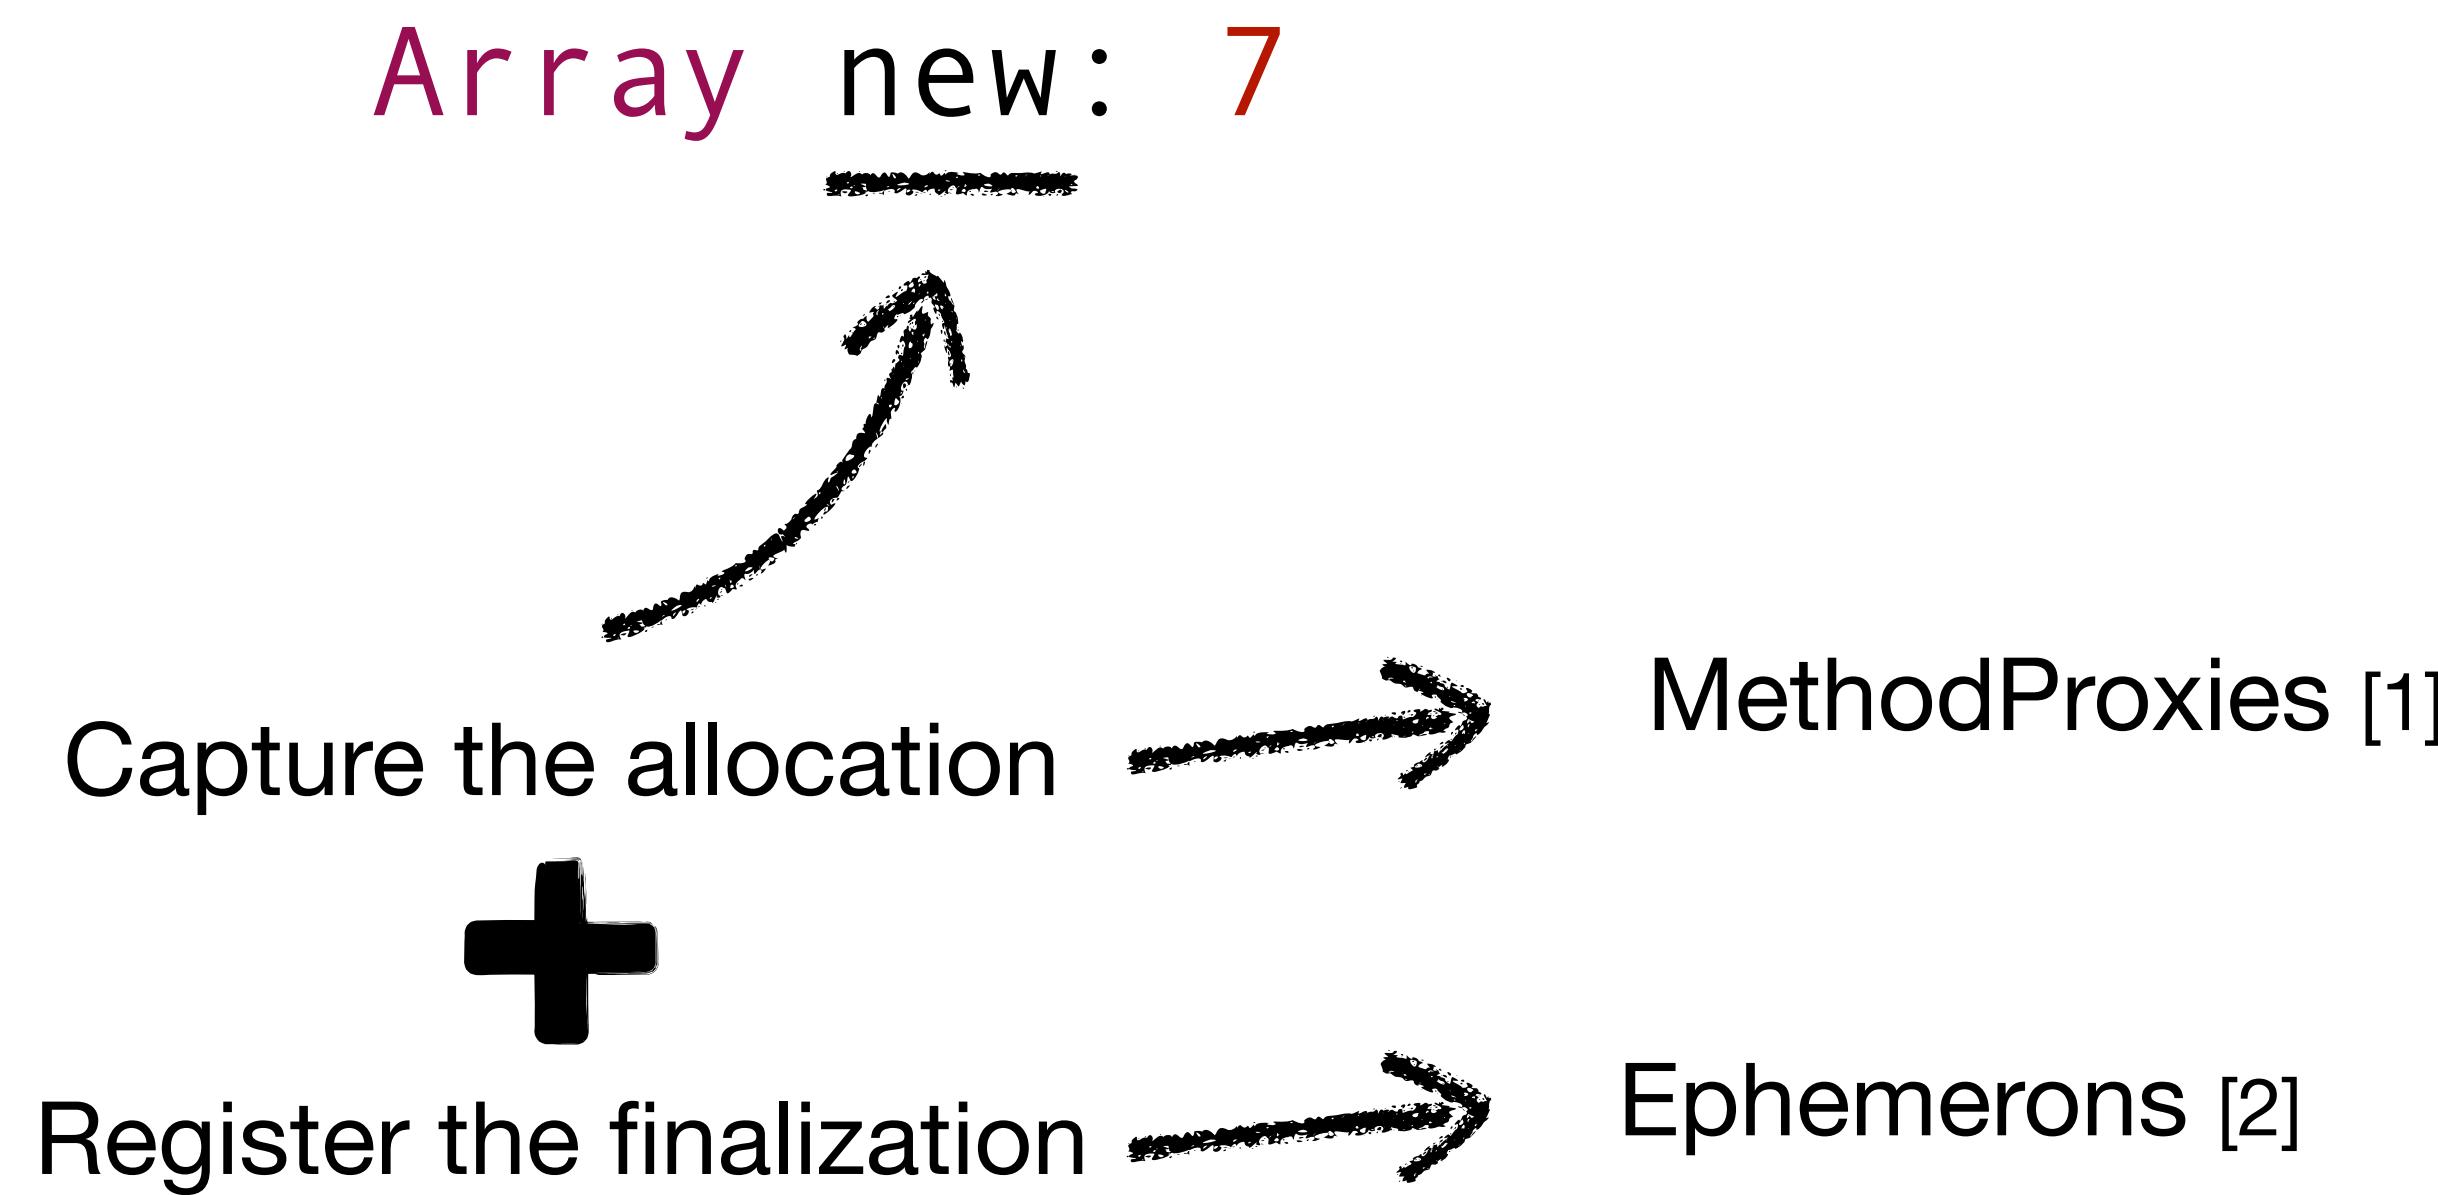
An object's lifetime



Capturing the allocations

Array new: 7

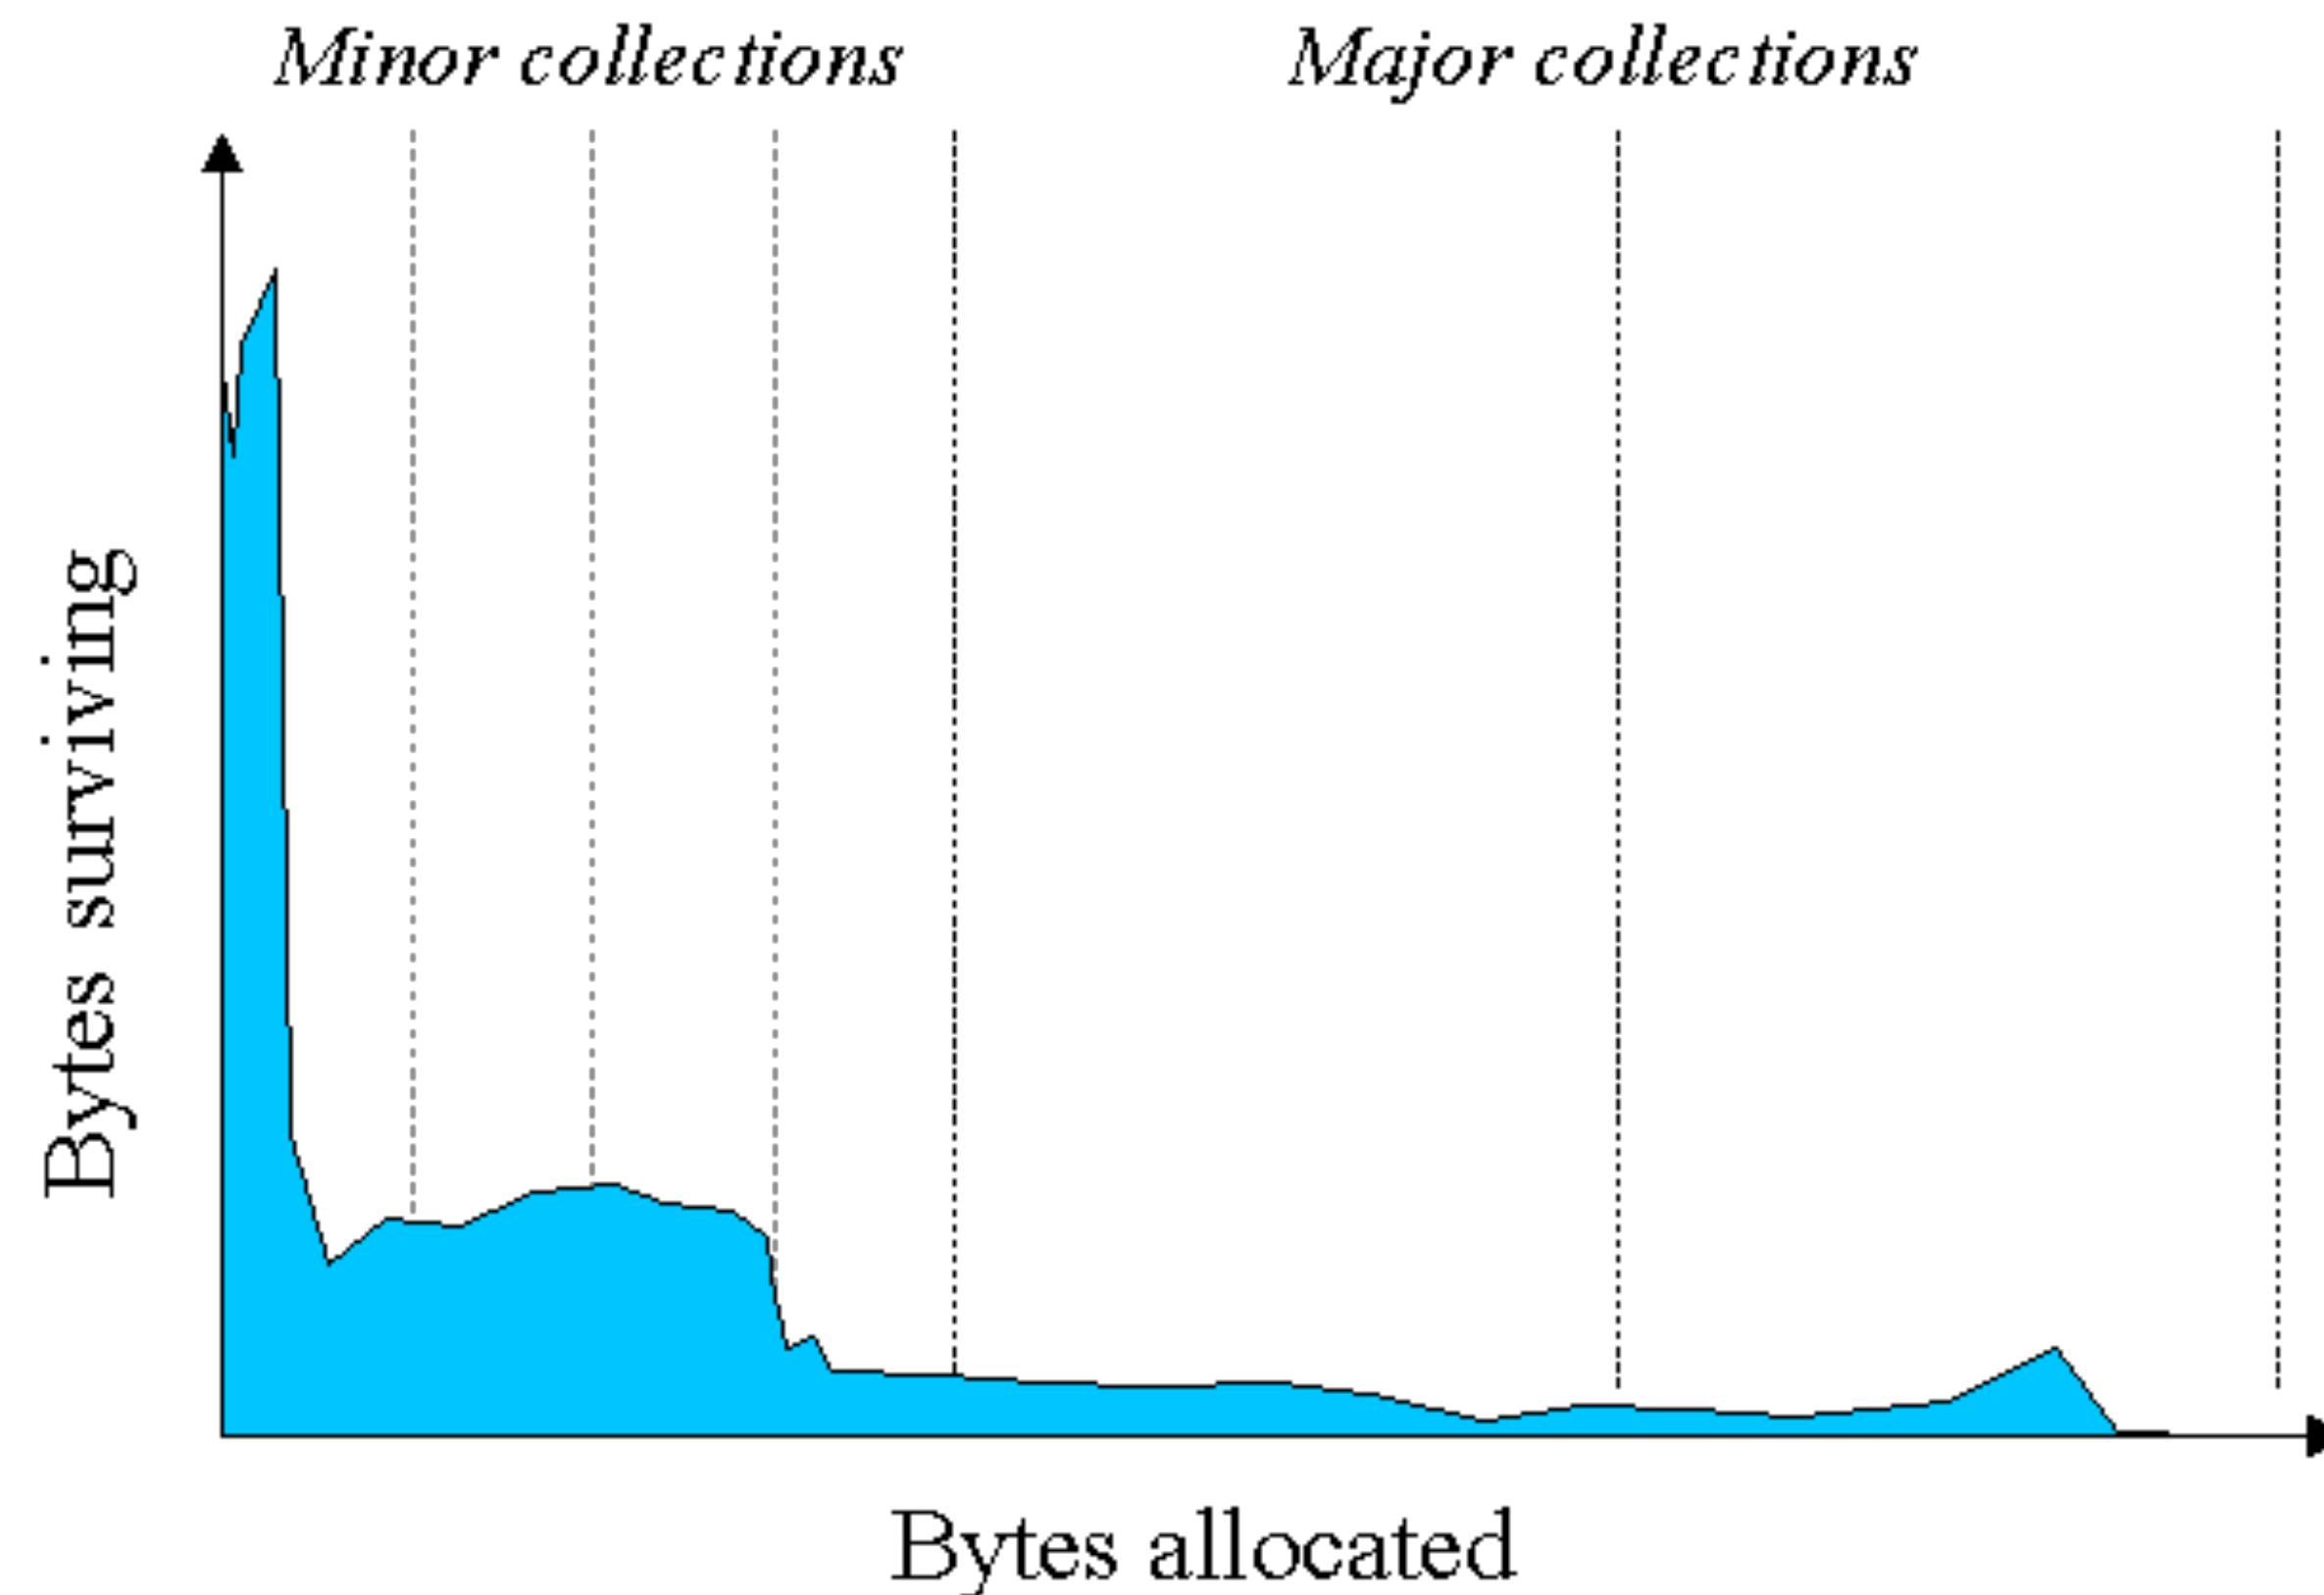
Capturing the allocations



[1] github.com/pharo-contributions/MethodProxies

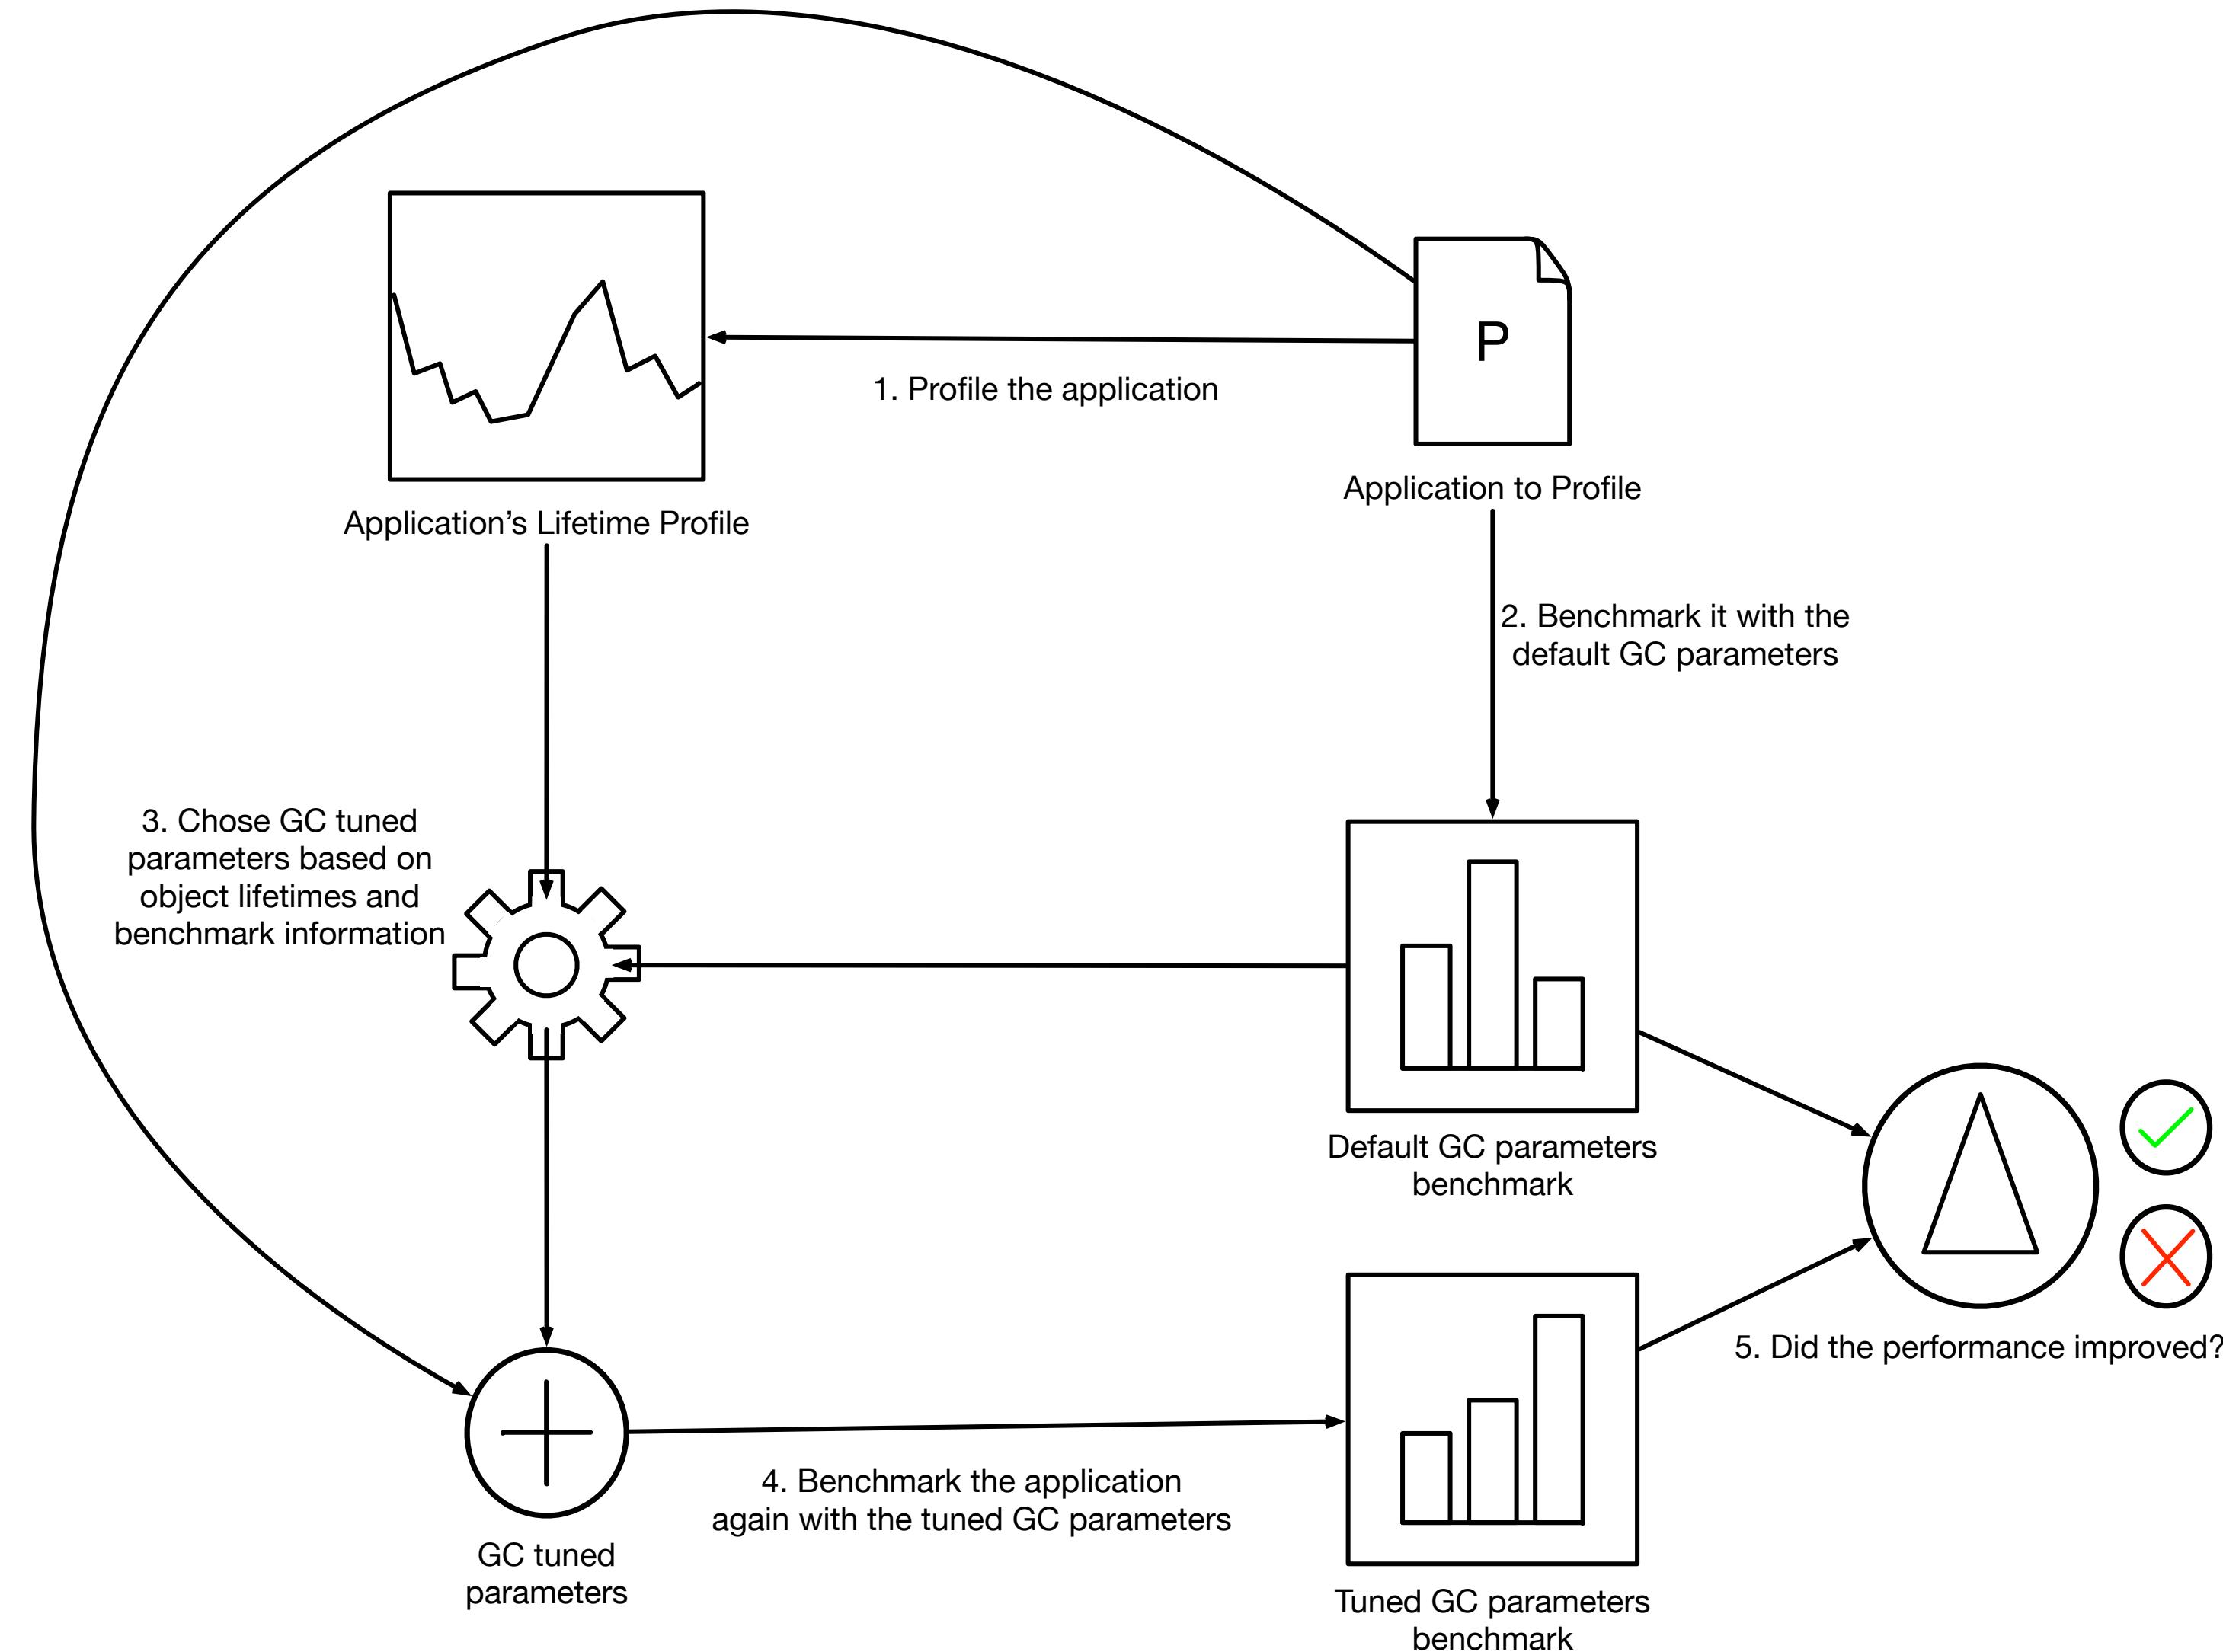
[2] github.com/pharo-project/pheps/blob/main/phep-0003.md

A common object lifetime distribution



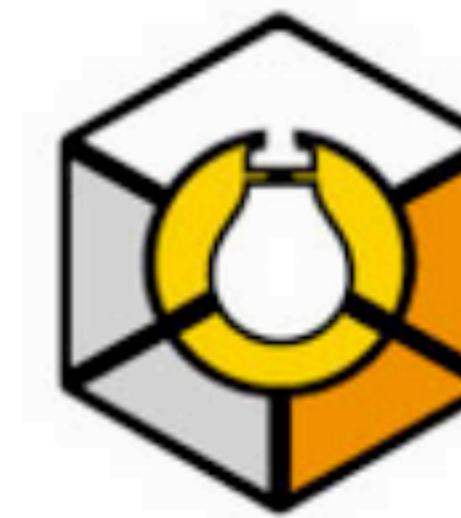
Source: oracle.com

Case study 2: methodology



Target application: DataFrame

PolyMathOrg/
DataFrame



DataFrame in Pharo - tabular data structures for
data analysis

12
Contributors

37
Issues

67
Stars

21
Forks



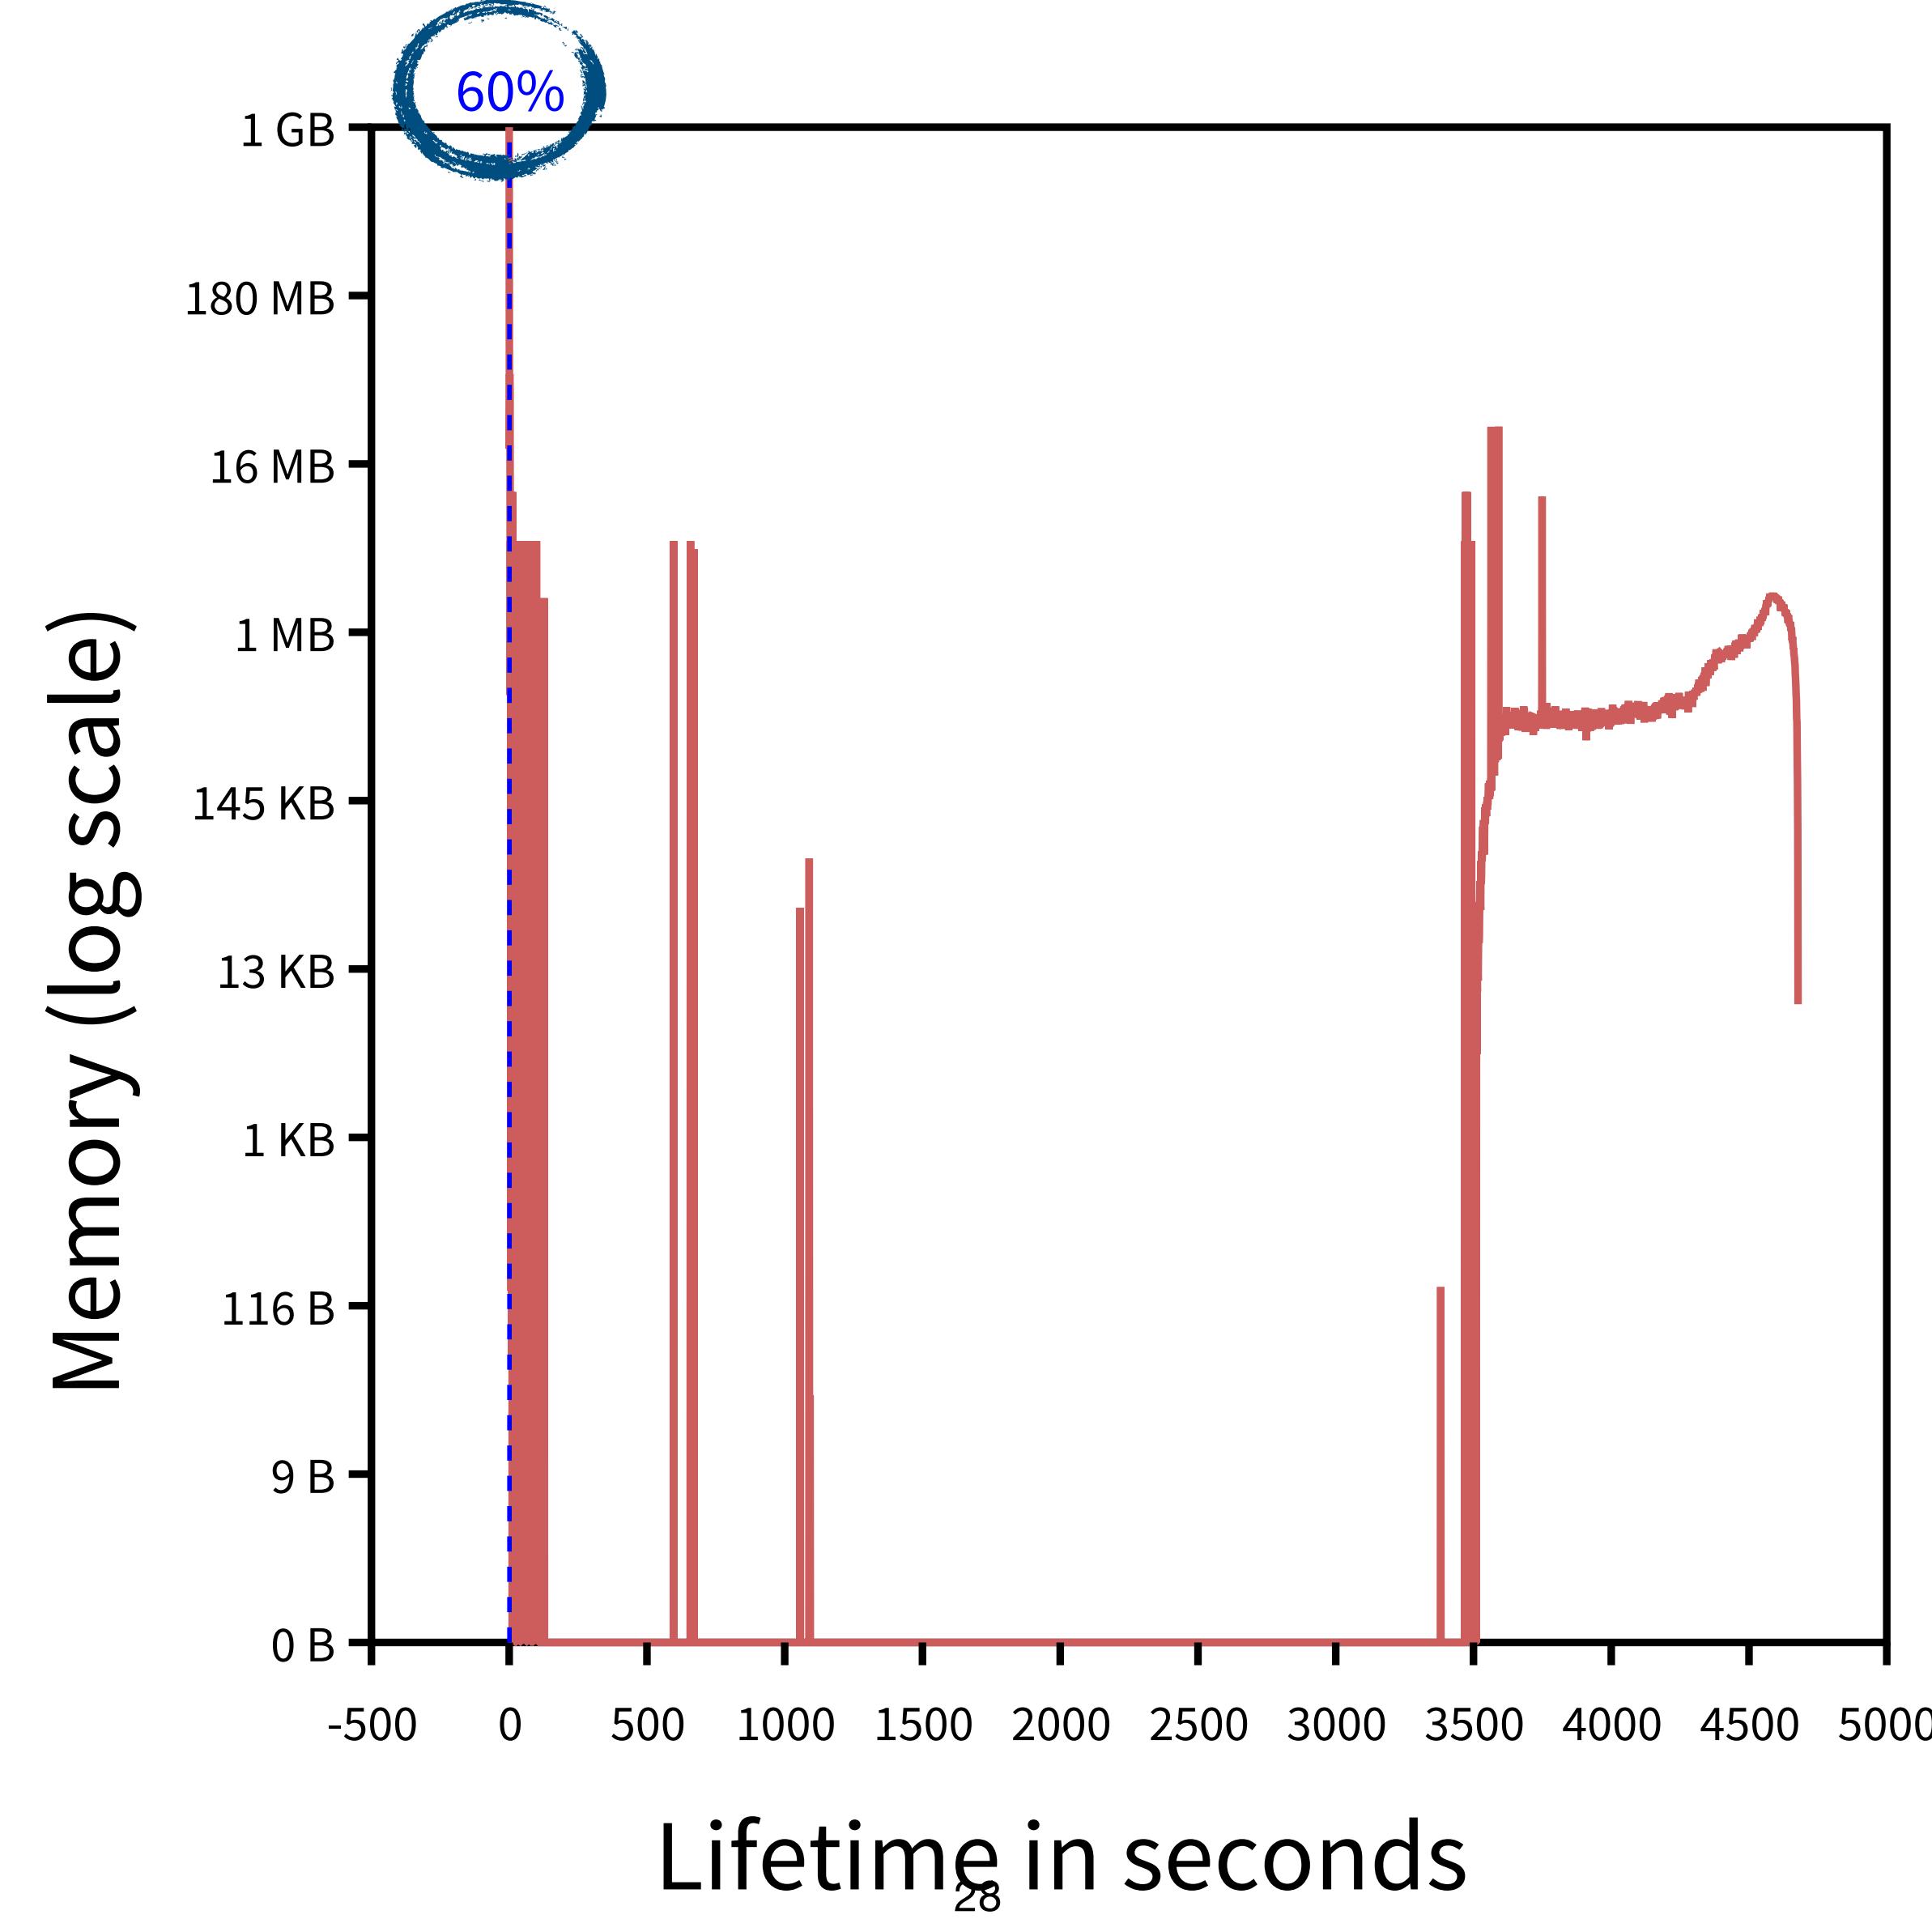
Benchmarking DataFrame

| Dataset | # of scavengers | # of full GCs | GC time | Total time | GC time in % |
|----------------|------------------------|----------------------|----------------|-------------------|---------------------|
| 500 MB | 266 | 18 | 11 sec | 71 sec | 15% |
| 1.6 GB | 304 | 36 | 60 sec | 248 sec | 22% |
| 3.1 GB | 1143 | 309 | 3793 sec | 4265 sec | 89% |

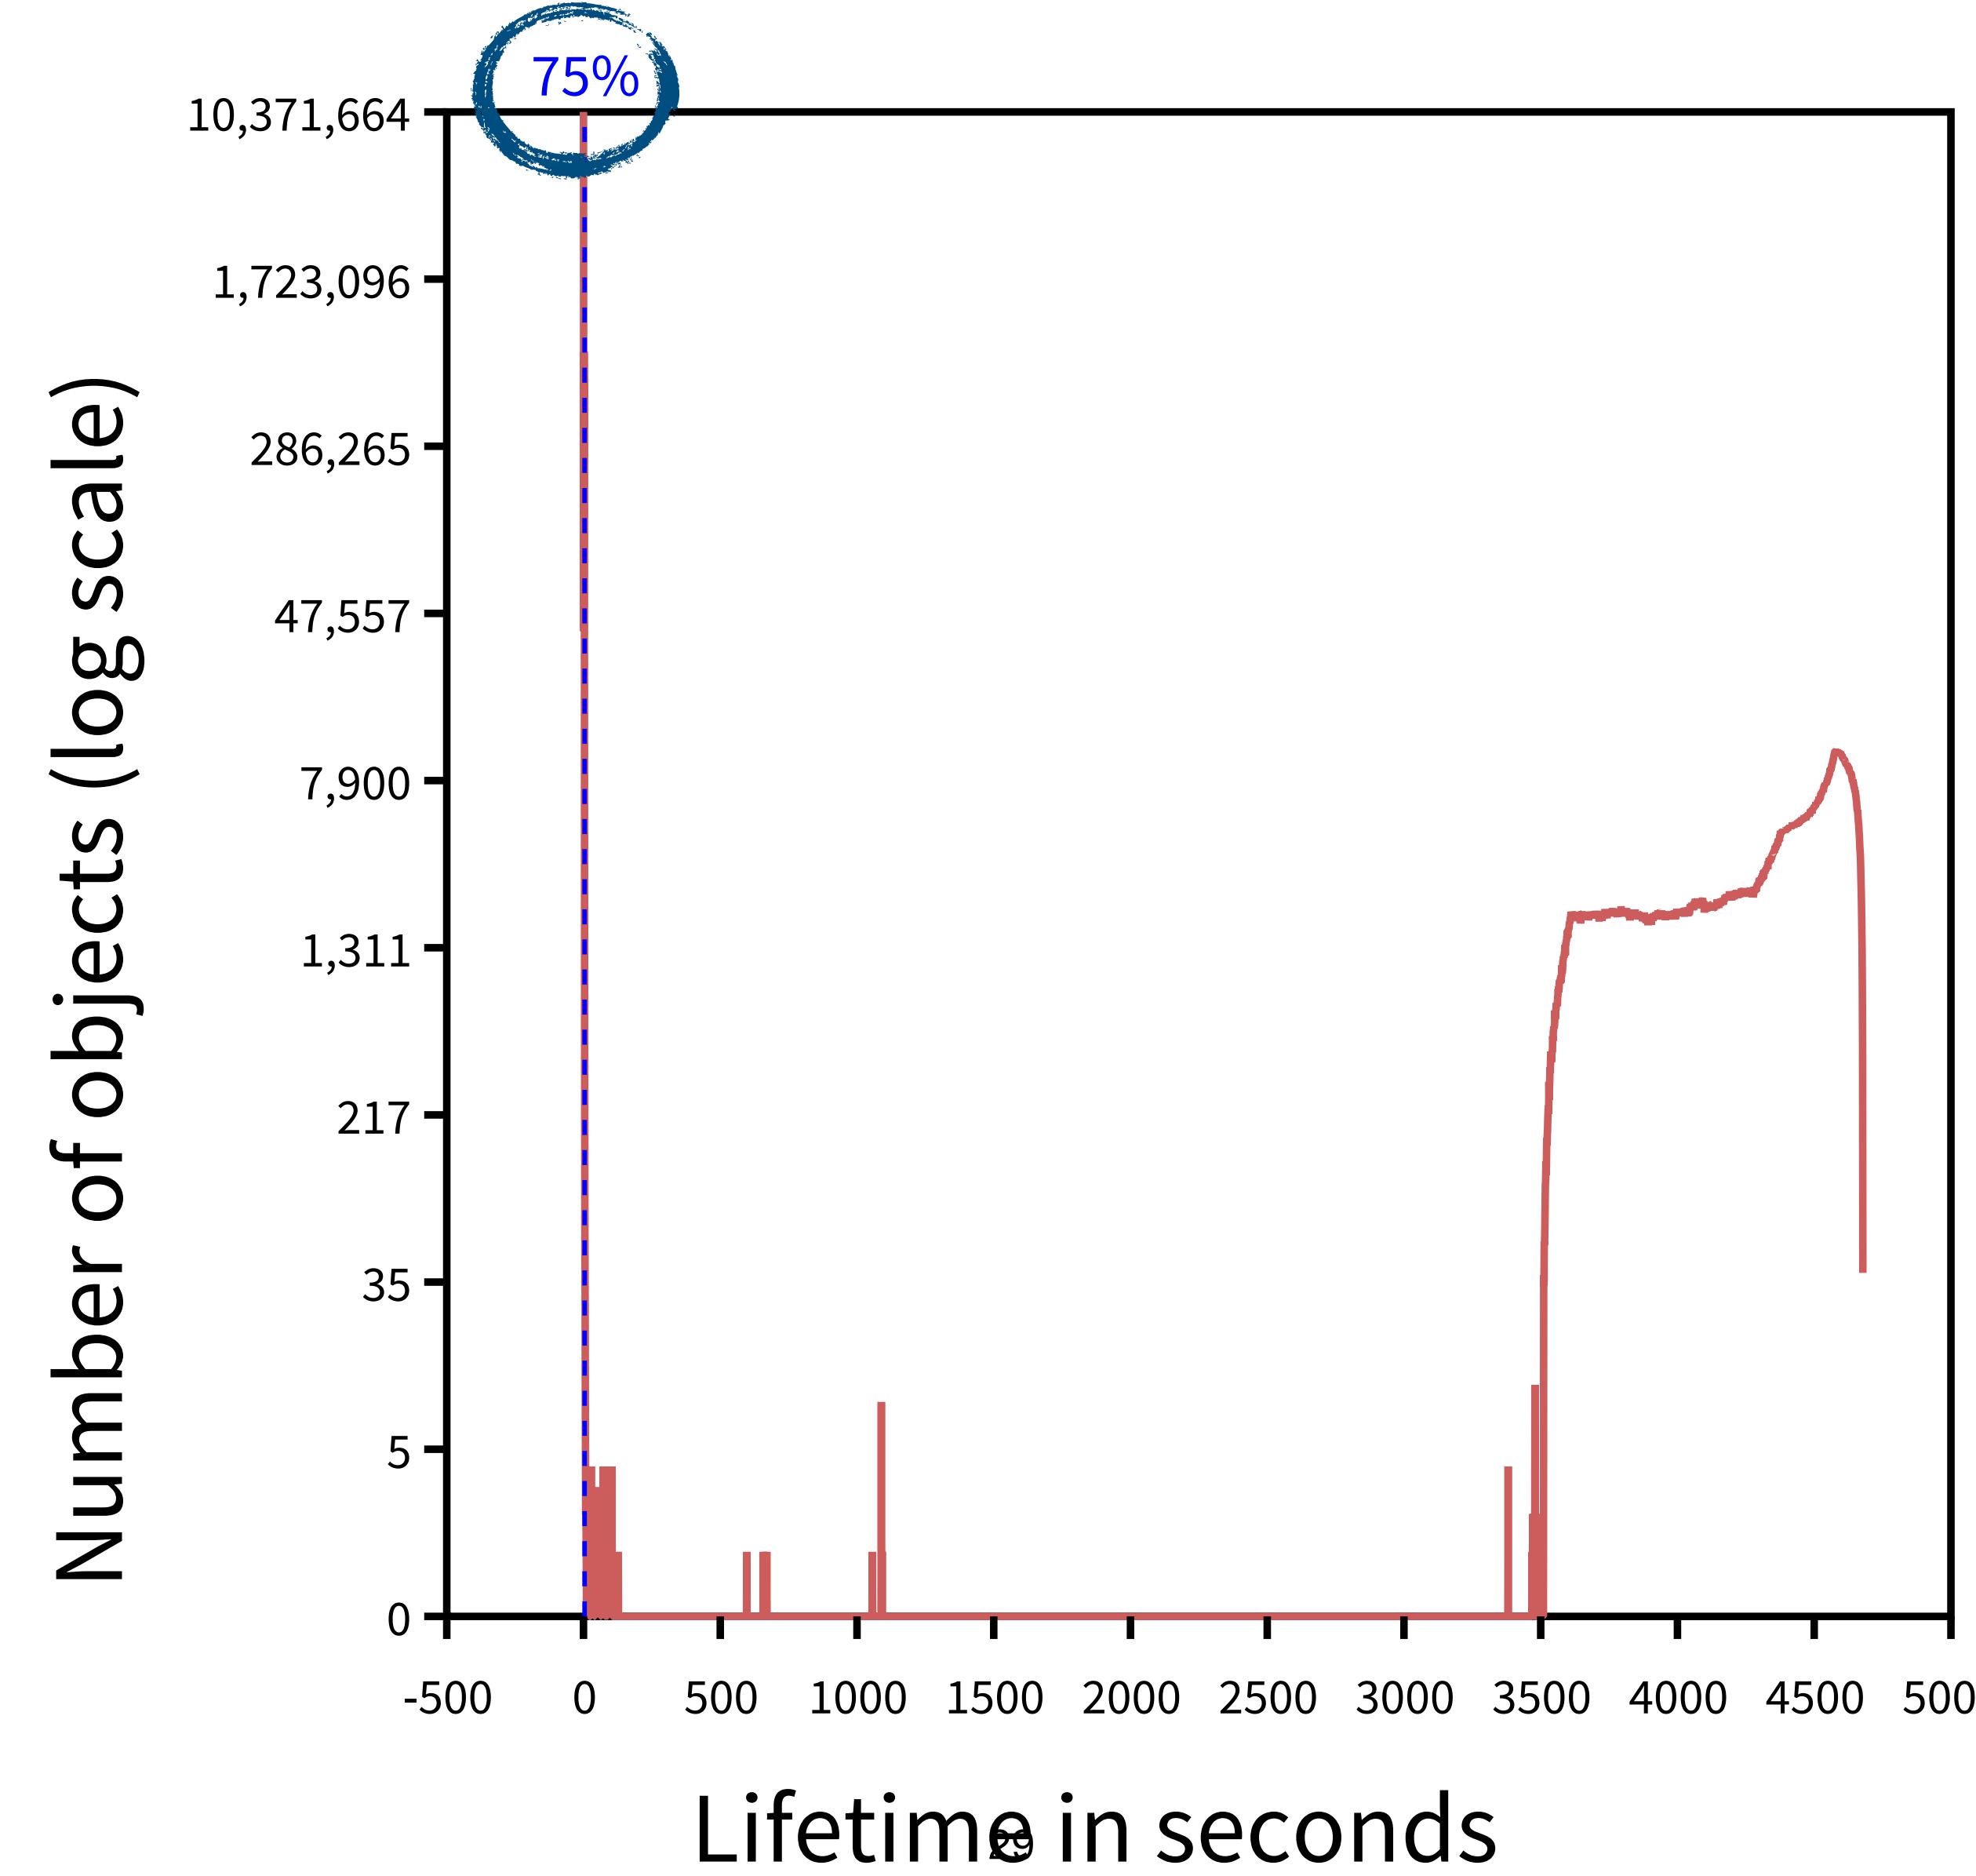
Benchmarking DataFrame

| Dataset | # of scavengers | # of full GCs | GC time | Total time | GC time in % |
|----------------|------------------------|----------------------|----------------|-------------------|---------------------|
| 500 MB | 266 | 18 | 11 sec | 71 sec | 15% |
| 1.6 GB | 304 | 36 | 60 sec | 248 sec | 22% |
| 3.1 GB | 1143 | 309 | 3793 sec | 4265 sec | 89% |

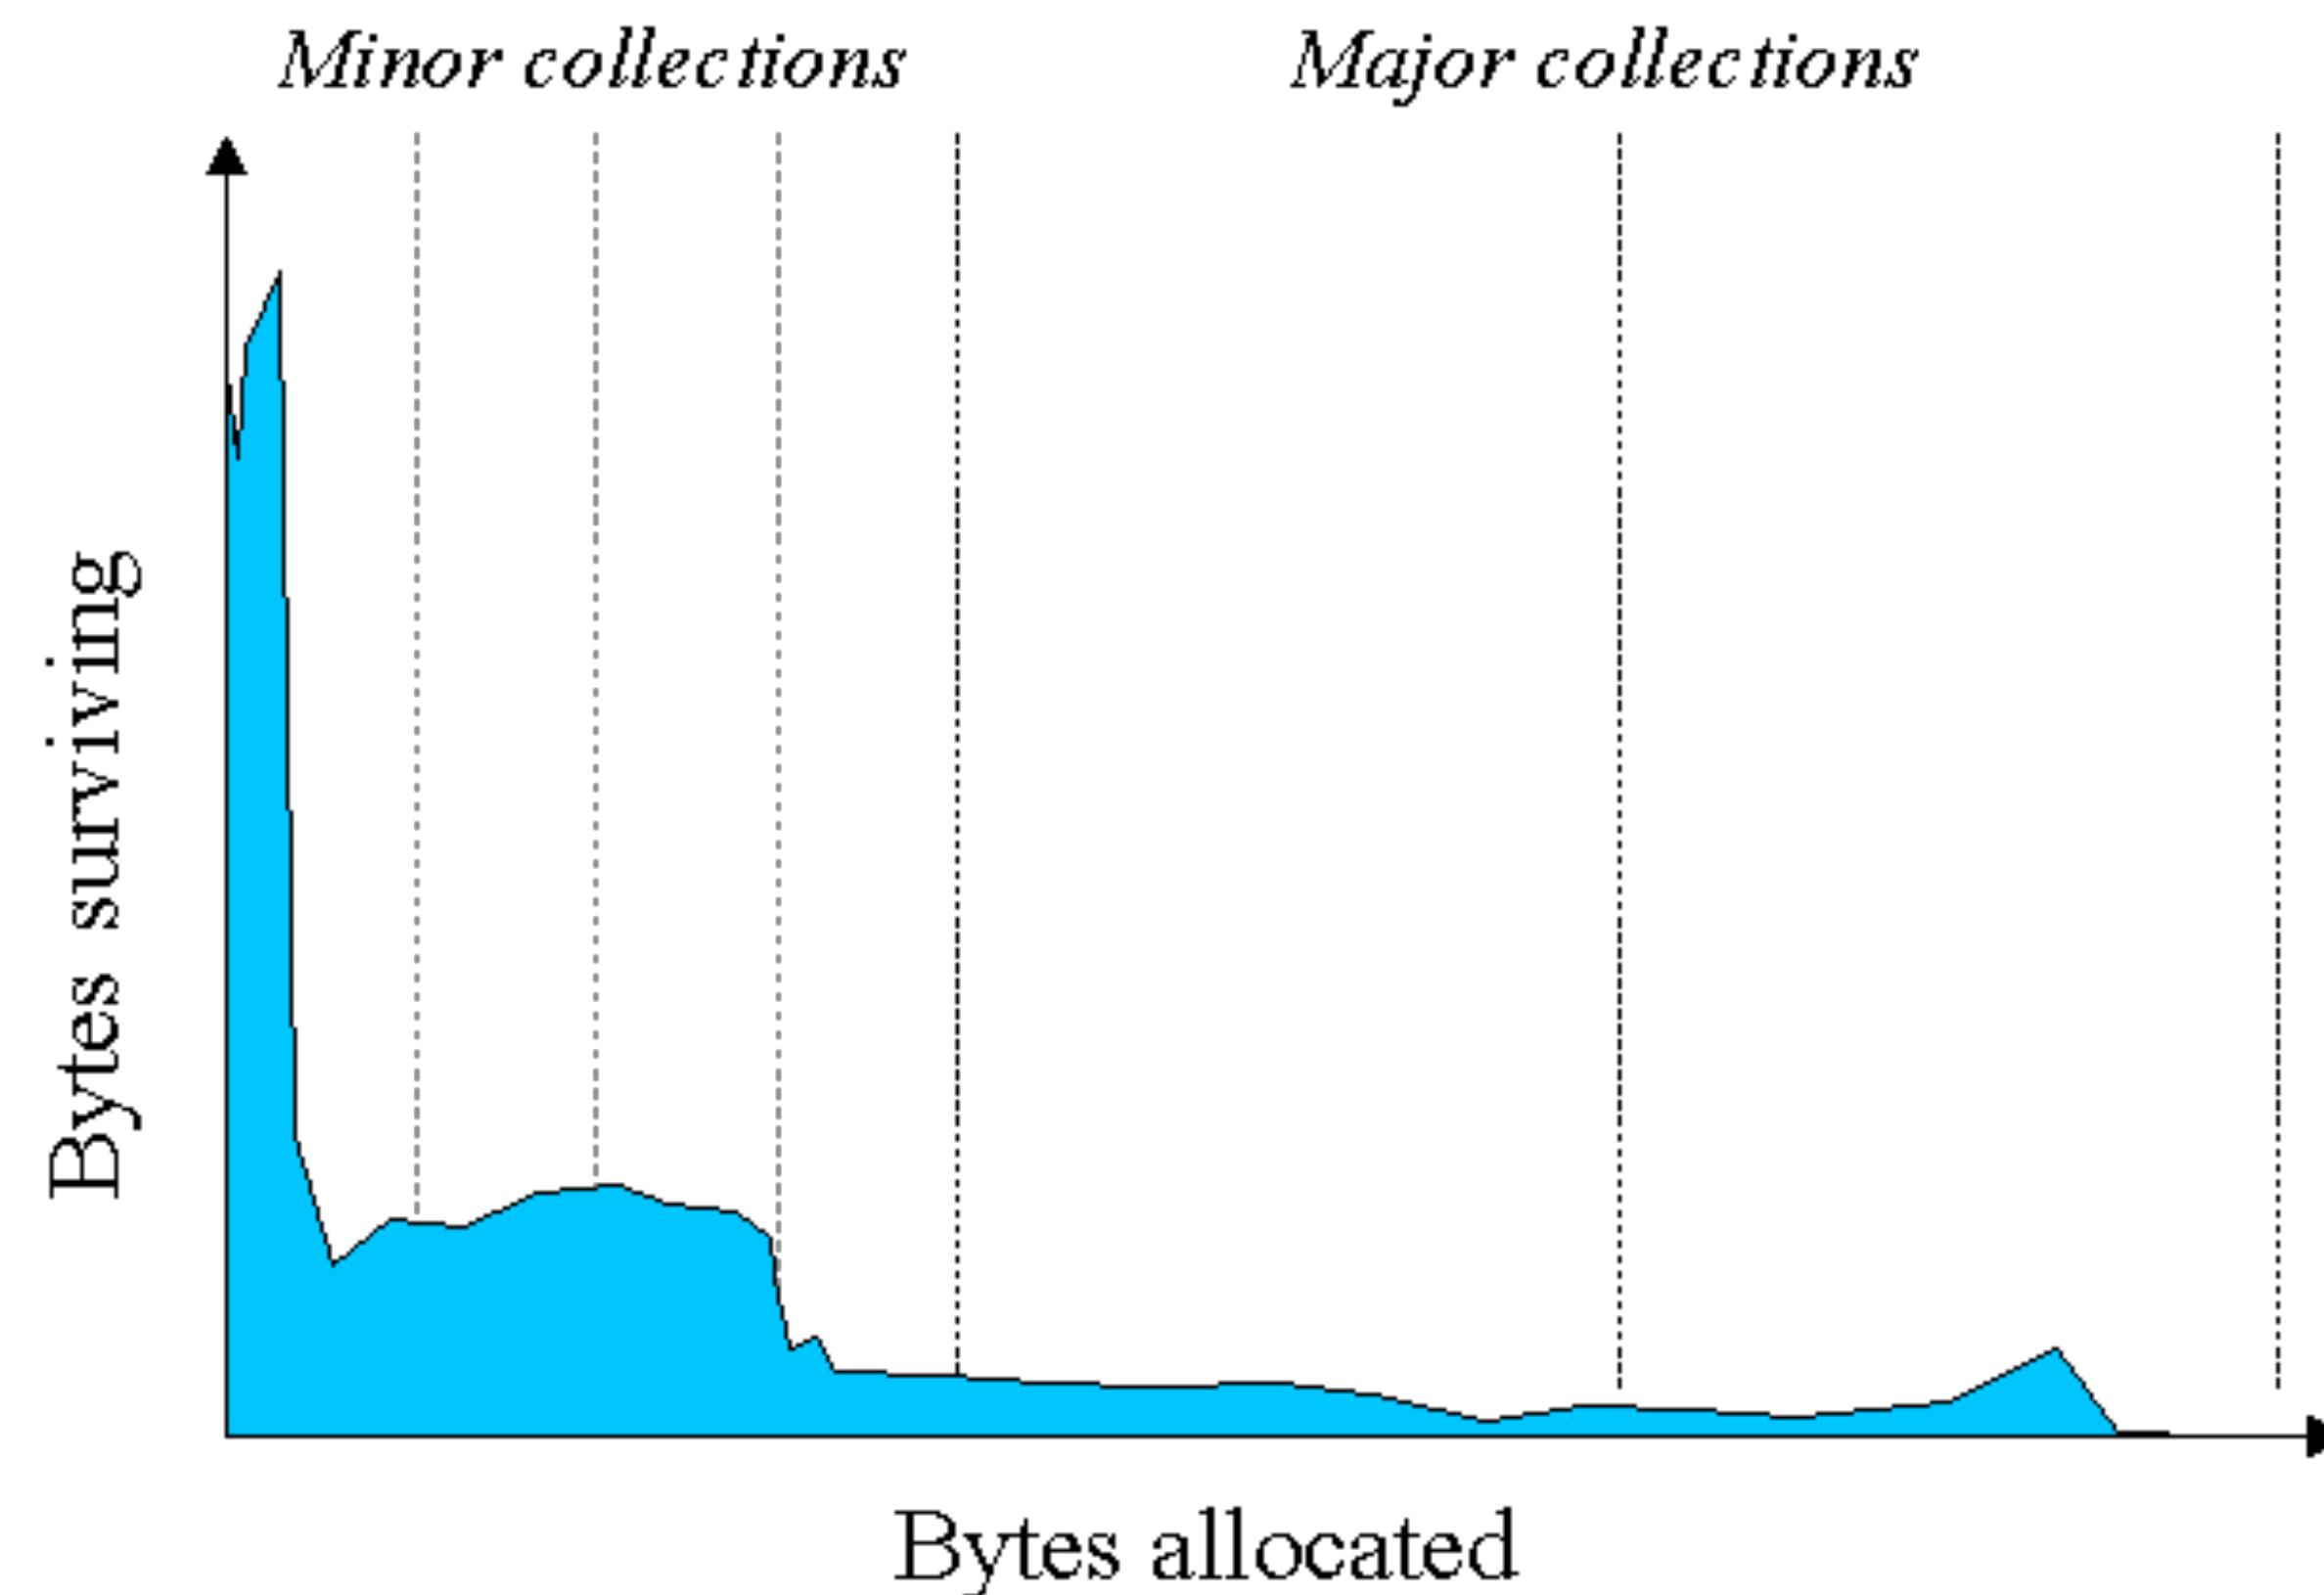
Object lifetimes for a 500MB DataFrame (memory)



Object lifetimes for a 500MB DataFrame (# objects)

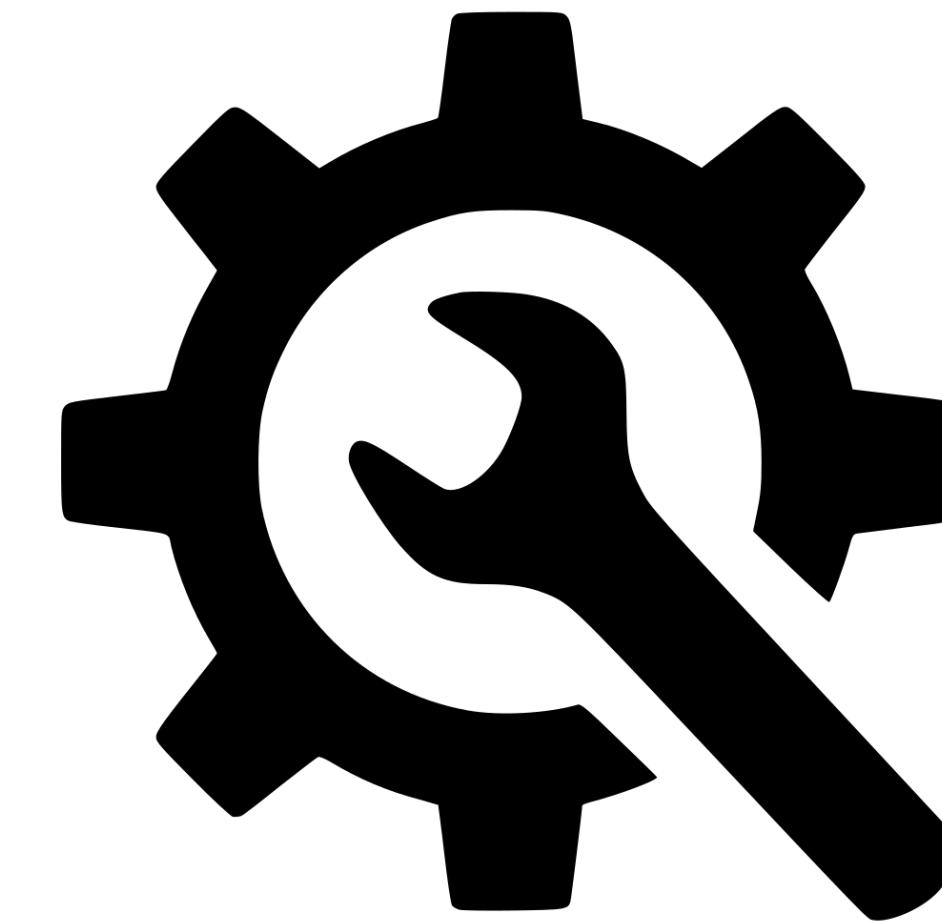
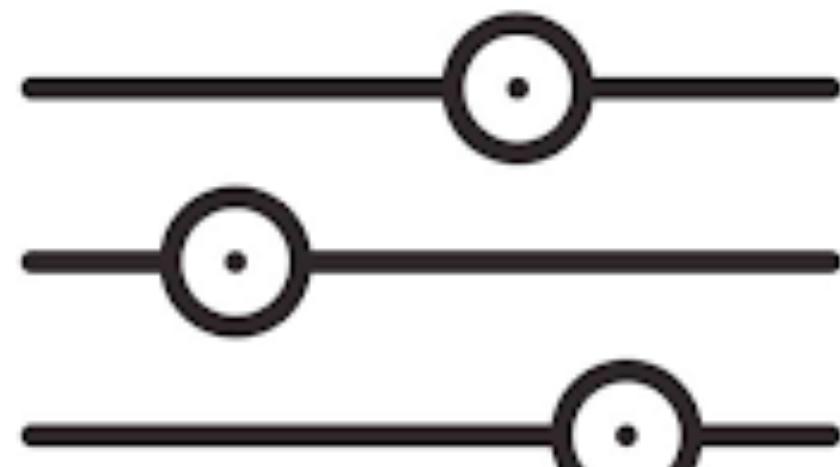


A common object lifetime distribution



Source: oracle.com

GC parameters



DataFrame performance improvements

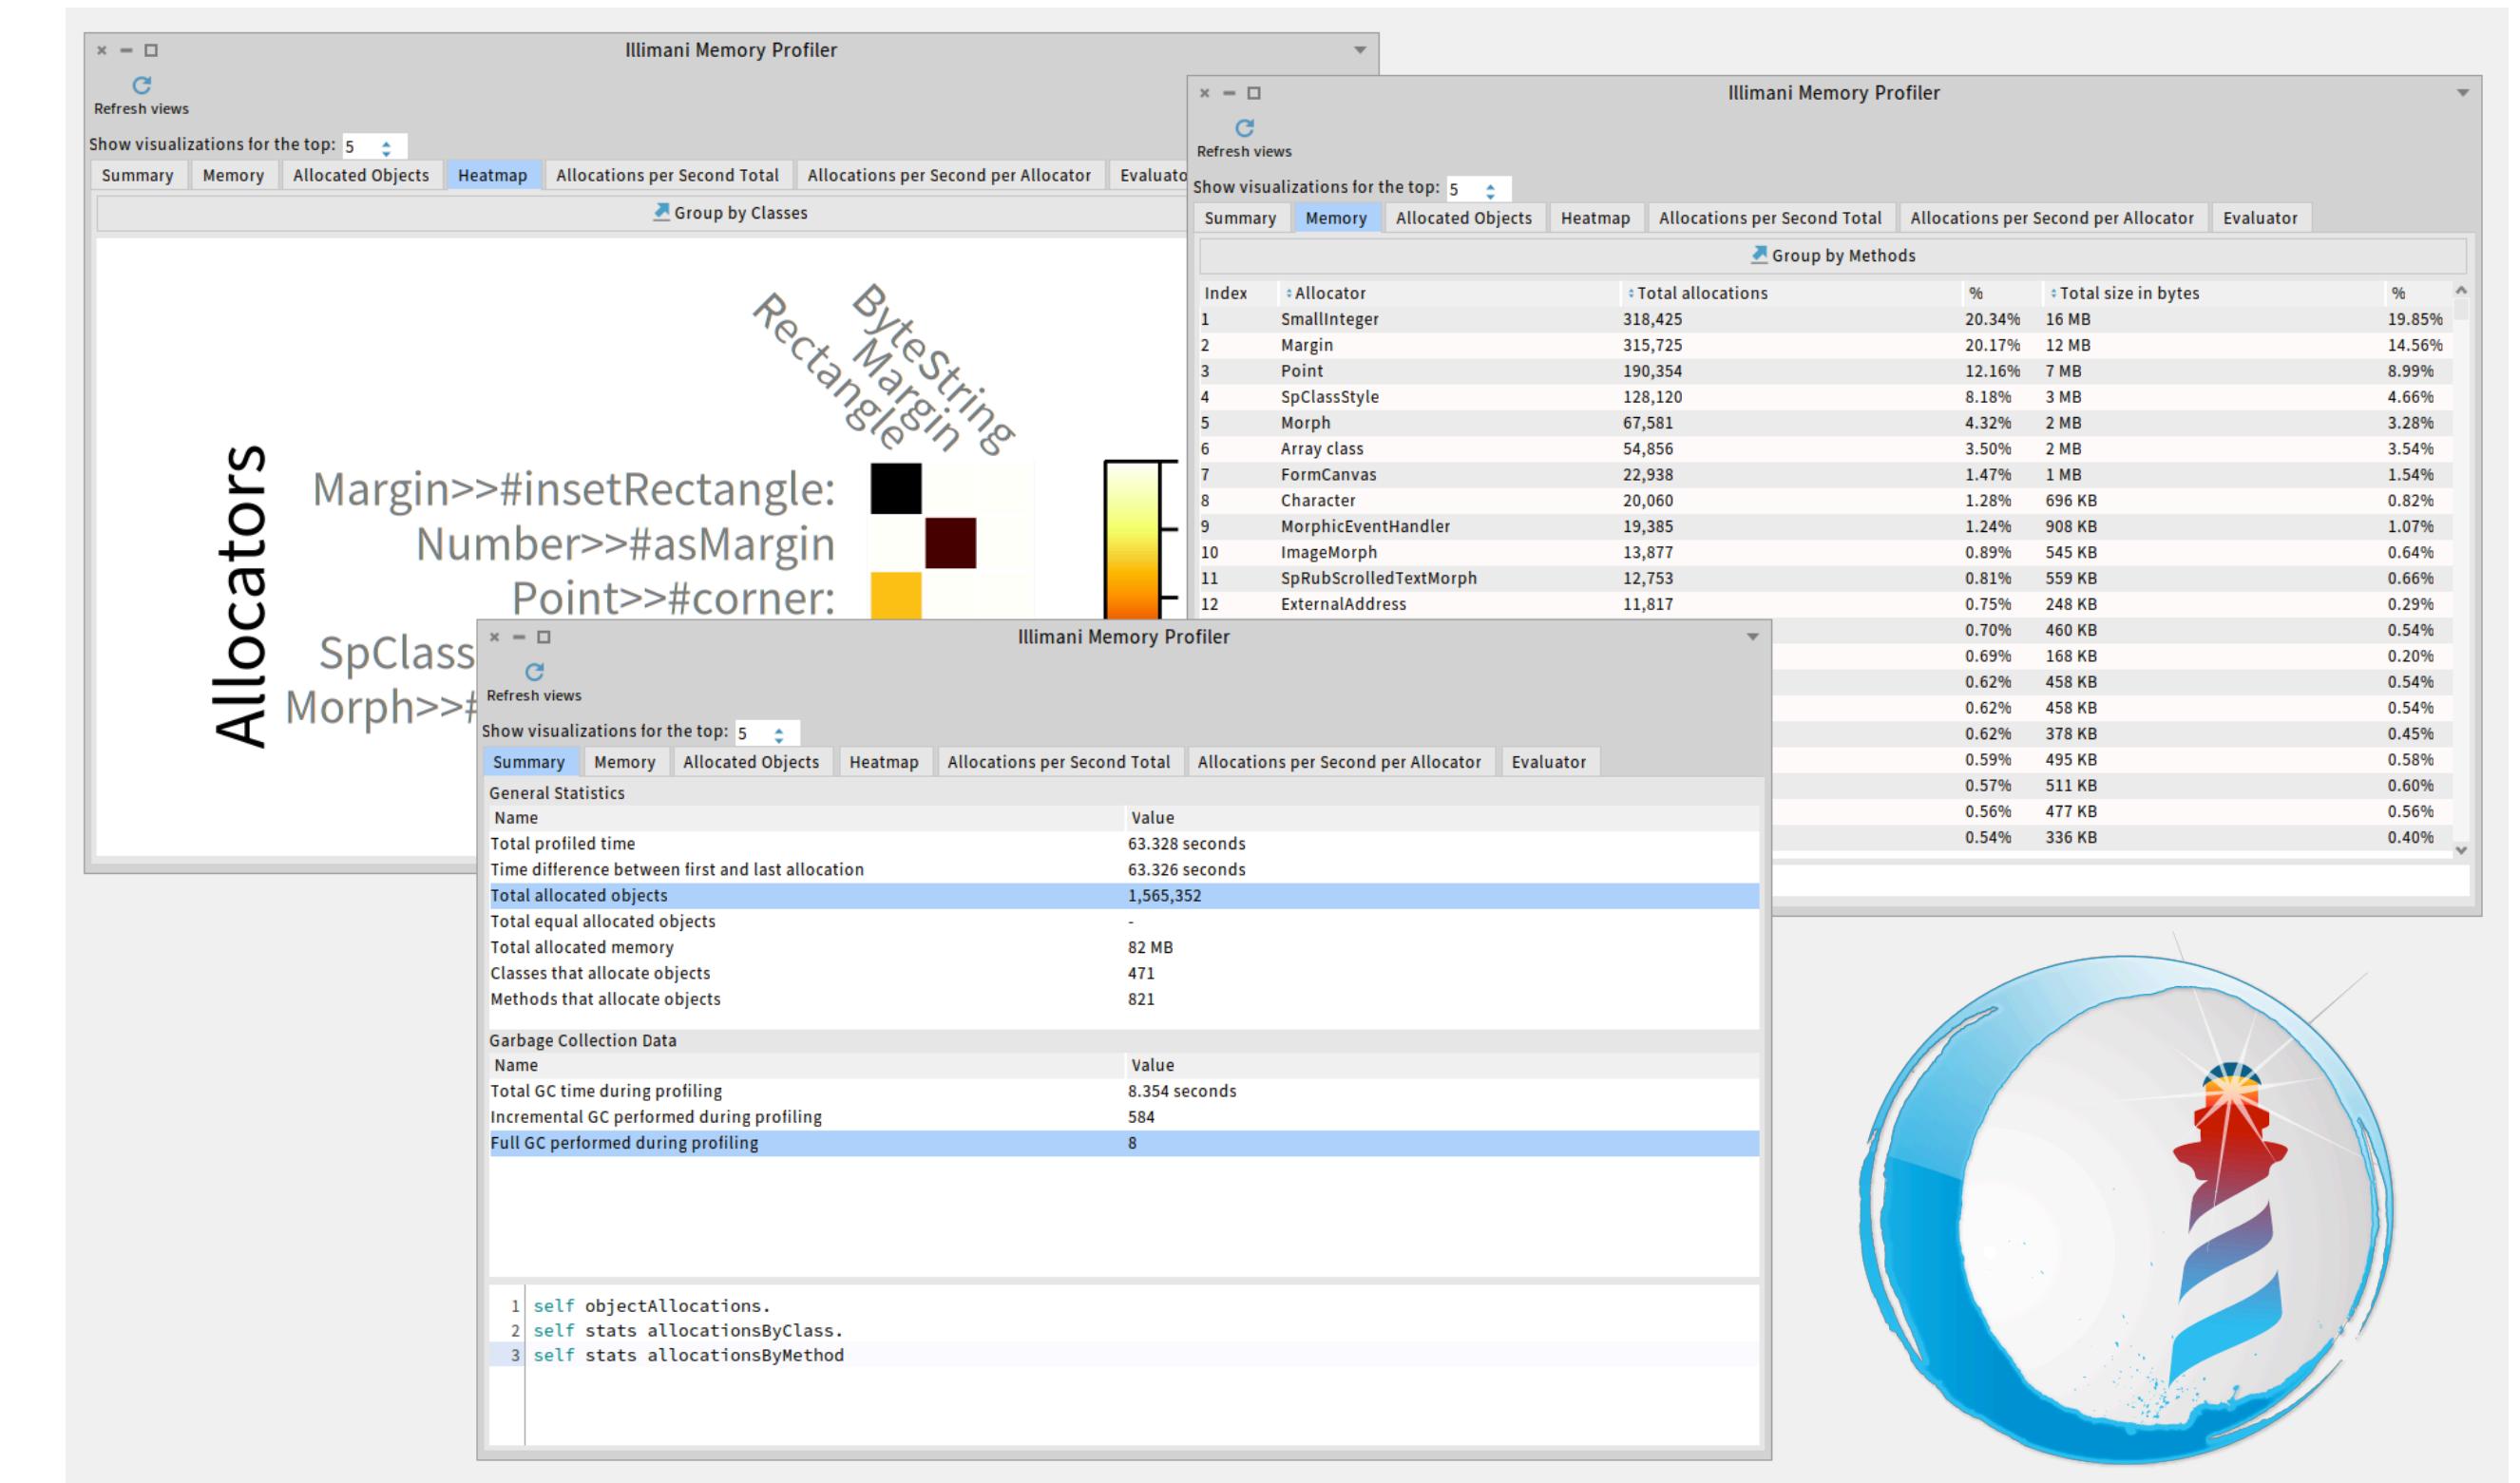
| GC Configuration | GC spent time | Total execution time | Improved performance |
|-------------------------|----------------------|-----------------------------|-----------------------------|
| Default | 58 min 18 sec | 1 hour 6 min 18 sec | 1× |
| Configuration 1 | 9 min 41 sec | 17 min 46 sec | 3.7× |
| Configuration 2 | 4 min 57 sec | 12 min 54 sec | 5.1× |
| Configuration 3 | 5 min 8 sec | 13 min 2 sec | 5.1× |
| Configuration 4 | 2 min 42 sec | 10 min 37 sec | 6.2× |
| Configuration 5 | 1 min 47 sec | 9 min 42 sec | 6.8× |

The future

- Study the precision of the approximated lifetimes.
- Dynamic optimizations based on allocation sites.

Illimani: a Pharo memory profiler

- Open-source MIT license
- Detects object allocation sites
- Tracks object lifetimes
- Allocation matrix
- Unmodified VM
- Density chart
- Memory consumption tables
- Rich object-oriented model



Sebastian JORDAN MONTAÑO
sebastian.jordan@inria.fr



github.com/jordanmontt/illimani-memory-profiler