

# Evaluating Finalization-Based Object Lifetime Profiling

**Sebastian JORDAN MONTAÑO**, Guillermo POLITO, Stéphane DUCASSE, Pablo TESONE

Inria, Univ. Lille, CNRS, Centrale Lille, UMR 9189 - CRISTAL

*sebastian.jordan@inria.fr*

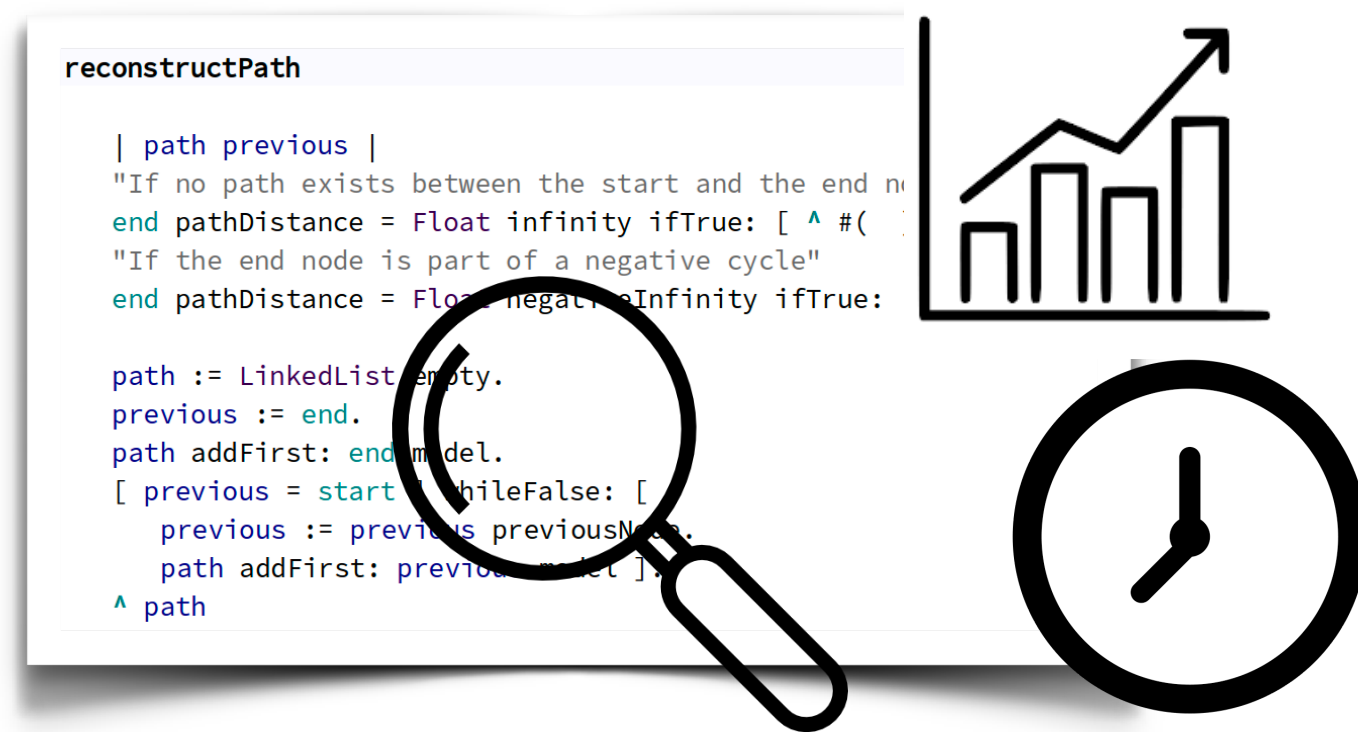
**International Symposium on Memory Management (ISMM) 2024**




**June 2024**

# Context and description

# Advantages of Object Lifetime Profiling



Object lifetime profiler



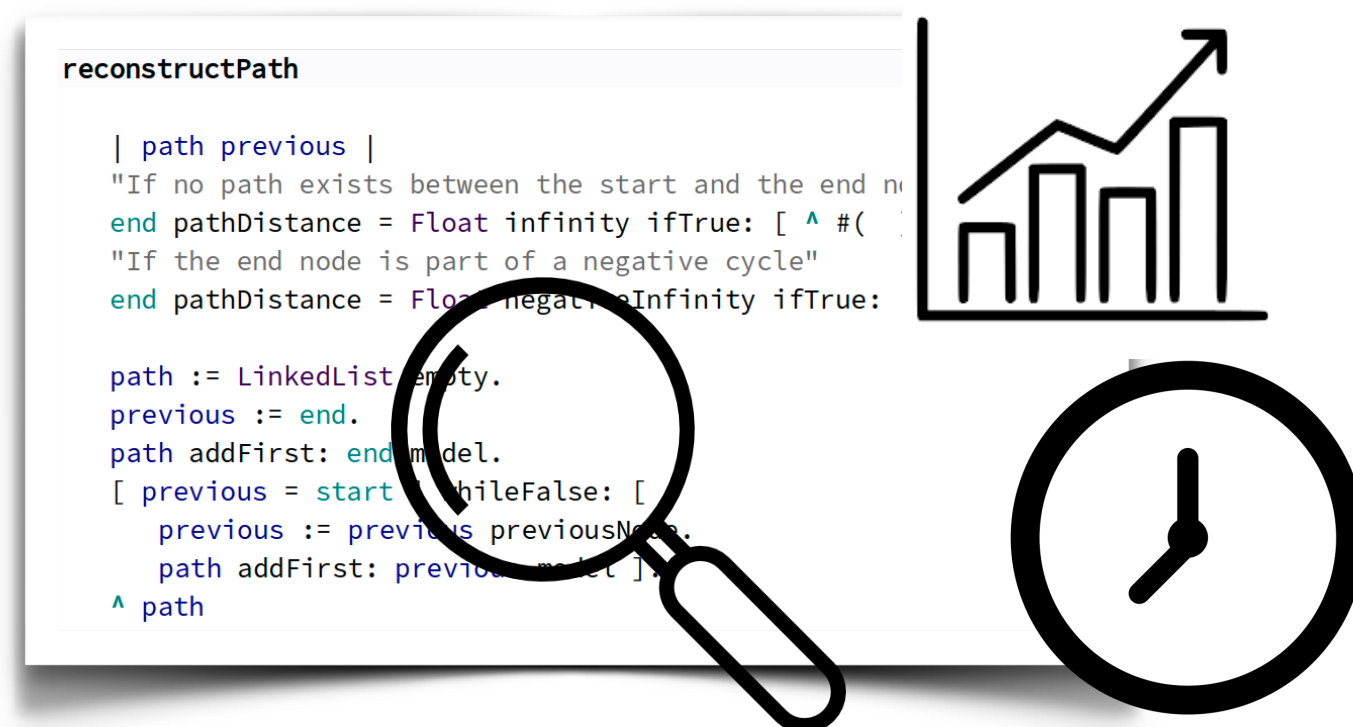
Useful for



Memory optimizations

- Pretenuring
- GC tuning

# Profiling Object Lifetimes with Finalization



Object lifetime profiler

It can be implemented using



Finalization mechanisms

# Finalization-Based Lifetime Profiling

$$\text{lifetime} = \underline{t_{\text{collection}}} - \underline{t_{\text{allocation}}}$$

Object Finalization


Code Instrumentation

The diagram illustrates the formula for lifetime profiling. The term  $t_{\text{collection}}$  is associated with 'Object Finalization' via a red line. The term  $t_{\text{allocation}}$  is associated with 'Code Instrumentation' via a red line. Both terms in the formula are underlined in red.

# Delayed Lifetime Observation

Approximated lifetime

Actual lifetime




**Lifetime observation time is delayed**

# Ephemeron Finalization and Contamination

- Finalization notifies object collection
- Ephemerons track objects of interest
- GC queues ephemeron when object is candidate for collection
- The execution engine dequeues and invokes finalization callback

1. GC queues object



2. Runtime invokes finalization callback

**Ephemeron contaminate the heap,  
adding extra GC pressure**

# Evaluating Finalization-Based Lifetime Profiling

- **Goal:** *understanding the impact of*
  - ephemeron contamination
  - delayed observation times



# FiLiP: **F**inalization **L**ifetime **P**rofiler

We built a profiler for experimentation: **FiLiP**

- Tracks allocation time by **instrumenting allocations**
- Tracks collection time with ***ephemeron finalization***
  - *optionally pre-tenured for experimentation*
- Built-in **sampling support**: track a *configurable* subset of all allocations

# Methodology and Results

# Research Questions

- RQ1: Are lifetime profilers actionable?
- **RQ2: How does sampling impact the computed object lifetimes?**
- **RQ3: Does sampling reduce the execution time overhead?**
- RQ4: Does sampling reduce the memory overhead?

# Benchmark selection

- **DataFrame**: Load a synthetic dataset of 230 MB .
- **HoneyGinger**: Run a smoothed-particle hydrodynamics simulator.
- **ReMobidyc**: Run a multi-agent simulator of wolves chasing and eating goats.
- **Bloc**: Render moving figures that simulate the flocking behavior of birds.
- **Moose**: Load a software database of 13521 classes and 48087 methods into the Moose meta-model.

# General methodology

- We profiled each benchmark using **4 different sampling rates**: 0.1%, 1%, 50%, and 100%
- Presented measured **lifetimes are relative** to total benchmark execution time.
  - To account for profiling overhead
- Execution time overhead ***with and without ephemeron pre-tenuring***

# RQ2: Precision Methodology

- We compare lifetimes across all sampling rates to assess variation.
- In-depth analysis of the DataFrame benchmark.

# RQ2: Precision

## Sampling and lifetimes

	100% Sampling	50% Sampling	1% Sampling	0.1% Sampling	Avg±stdev
<b>DataFrame</b>					
Overall average lifetimes	15.24%	14.68%	14.09%	14.21%	14.55%±0.45
<b>HoneyGinger</b>					
Overall average lifetimes	0.09%	0.08%	0.05%	0.06%	0.07%±0.014
<b>Re:Mobidyc</b>					
Overall average lifetimes	0.5%	0.36%	0.46%	0.18%	0.37%±0.12
<b>Bloc</b>					
Overall average lifetimes	6.27%	6.27%	6.51%	6.63%	6.42%±0.15
<b>Moose</b>					
Overall average lifetimes	13.19%	13.33%	13.67%	12.82%	13.25%±0.3

# RQ2: Precision

## Sampling and lifetimes

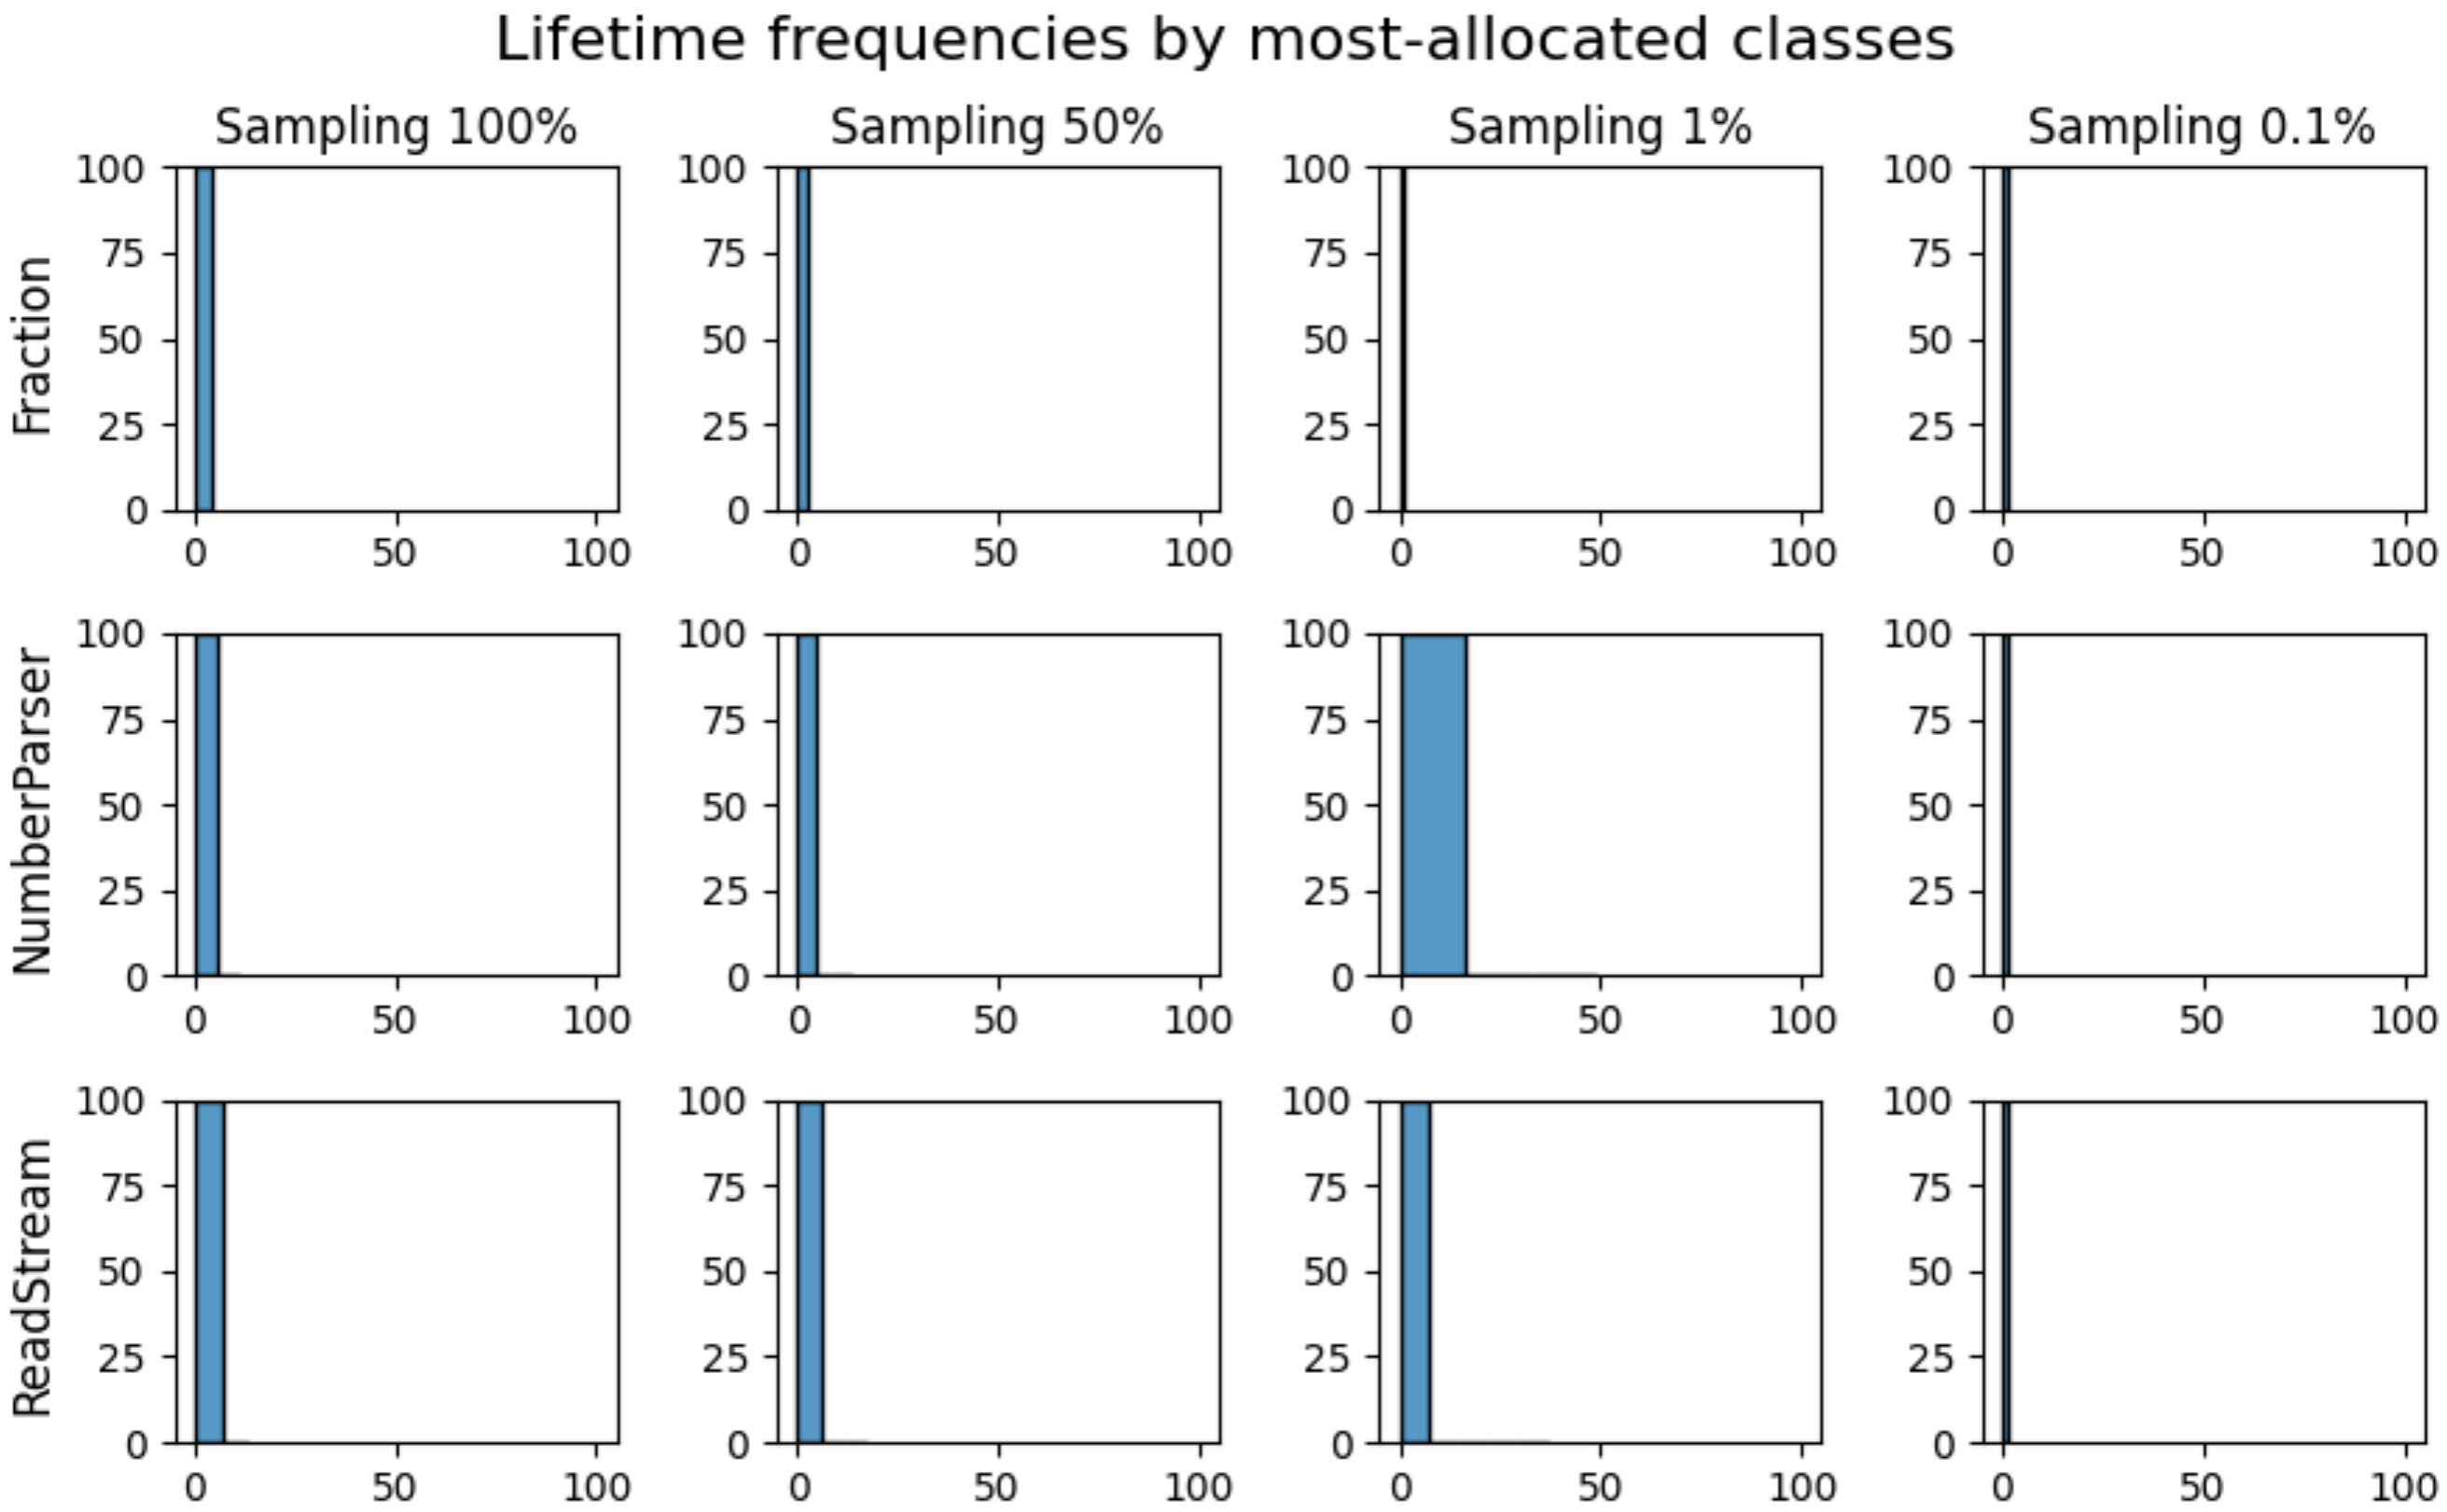
	100% Sampling	50% Sampling	1% Sampling	0.1% Sampling	Avg±stdev
<b>DataFrame</b>					
Overall average lifetimes	15.24%	14.68%	14.09%	14.21%	14.55%±0.45
<b>HoneyGinger</b>					
Overall average lifetimes	0.09%	0.08%	0.05%	0.06%	0.07%±0.014
<b>Re:Mobidyc</b>					
Overall average lifetimes	0.5%	0.36%	0.46%	0.18%	0.37%±0.12
<b>Bloc</b>					
Overall average lifetimes	6.27%	6.27%	6.51%	6.63%	6.42%±0.15
<b>Moose</b>					
Overall average lifetimes	13.19%	13.33%	13.67%	12.82%	13.25%±0.3




# RQ2: Precision

- DataFrame case
- Segregating short and long lived classes

Short-Lived



Long-Lived



# RQ2: Conclusion

- Consistent lifetime frequencies **suggest precise measurements.**
- **Recommendation for developers:**
  - Start with a small sampling rate
  - Gradually incrementing it if needed
  - We obtained **precise results using a 0.1% sampling rate**

# RQ3: Execution Time Overhead Methodology

- Comparison across **5 configurations**:
  - Baseline: without the profiler
  - Profiler active with sampling rates of 100%, 50%, 1% and 0.1%.
- Effect of **pretenuring the ephemerons**

# RQ3: Execution time overhead

Sampling rate	100%	50%	1%	0.1%
DataFrame - pretenuring	$5.8 \times \pm 0.017$	$3.8 \times \pm 0.019$	$2.25 \times \pm 0.019$	$2.27 \times \pm 0.012$
DataFrame + pretenuring	$2.29 \times \pm 0.03$	$1.64 \times \pm 0.01$	$1.08 \times \pm 0.001$	$1.12 \times \pm 0.001$
$\Delta$ Pretenuring	2.53×	2.32×	2.08×	2.03×
HoneyGinger - pretenuring	$9.2 \times \pm 0.01$	$8.91 \times \pm 0.03$	$8.4 \times \pm 0.6$	$7.8 \times \pm 0.04$
HoneyGinger + pretenuring	$3.2 \times \pm 0.10$	$2.5 \times \pm 0.10$	$2.2 \times \pm 0.06$	$1.8 \times \pm 0.05$
$\Delta$ Pretenuring	2.88×	3.56×	3.82×	4.33×
Bloc - pretenuring	$1.04 \times \pm 0.002$	$1.04 \times \pm 0.001$	$1.04 \times \pm 0.0008$	$1.04 \times \pm 0.0006$
Bloc + pretenuring	$1.03 \times \pm 0.001$	$1.03 \times \pm 0.001$	$1.03 \times \pm 0.0007$	$1.03 \times \pm 0.0007$
$\Delta$ Pretenuring	1.01×	1.01×	1.01×	1.01×
Moose - pretenuring	$3.2 \times \pm 0.02$	$2.6 \times \pm 0.008$	$2.02 \times \pm 0.007$	$2.00 \times \pm 0.005$
Moose + pretenuring	$2.02 \times \pm 0.03$	$1.6 \times \pm 0.007$	$1.1 \times \pm 0.003$	$1.09 \times \pm 0.002$
$\Delta$ Pretenuring	1.58×	1.63×	1.84×	1.84×
Re:Mobidyc - pretenuring	$32.3 \times \pm 2.2$	$13.8 \times \pm 0.11$	$9.14 \times \pm 0.13$	$9.01 \times \pm 0.08$
Re:Mobidyc + pretenuring	$2.06 \times \pm 0.005$	$1.5 \times \pm 0.004$	$1.3 \times \pm 0.01$	$1.25 \times \pm 0.05$
$\Delta$ Pretenuring	15.68×	9.20×	7.03×	7.21×
Avg. + pretenuring	2.12×	1.65×	1.34×	1.26×

# RQ3: Conclusion

- Pretenuring ephemeron
  - **reduces the overhead by 3.68x** on average
  - **average overhead of 1.59x** across all sampling rates
- Finalization profiling is a **practical alternative to estimating object lifetimes**



# More in the paper!

## Evaluating Finalization-Based Object Lifetime Profiling

Sebastian Jordan Montaña

Univ. Lille - Inria - CNRS - Centrale Lille - UMR 9189  
CRISTAL  
Villeneuve D'Ascq, France  
sebastian.jordan@inria.fr

Stephane Ducasse

Univ. Lille - Inria - CNRS - Centrale Lille - UMR 9189  
CRISTAL  
Villeneuve D'Ascq, France  
stephane.ducasse@inria.fr

Guillermo Polito

Univ. Lille - Inria - CNRS - Centrale Lille - UMR 9189  
CRISTAL  
Villeneuve D'Ascq, France  
guillermo.polito@inria.fr

Pablo Tesone

Univ. Lille - Inria - CNRS - Centrale Lille - UMR 9189  
CRISTAL  
Villeneuve D'Ascq, France  
pablo.tesone@inria.fr

### Abstract

Using object lifetime information enables performance improvement through memory optimizations such as pretenuring and tuning garbage collector parameters. However, profiling object lifetimes is nontrivial and often requires a specialized virtual machine to instrument object allocations and dereferences. Alternative lifetime profiling could be done with less implementation effort using available finalization mechanisms such as weak references.

In this paper, we study the impact of finalization on object lifetime profiling. We built an actionable lifetime profiler using the ephemeron finalization mechanism named FiLiP. FiLiP instruments object allocations to exactly record an object's allocation time and it attaches an ephemeron to each allocated object to capture its finalization time. We show that FiLiP can be used in practice and achieves a significant overhead reduction by pretenuring the ephemeron objects. We further experiment with the impact of sampling allocations, showing that sampling reduces profiling overhead while maintaining actionable lifetime measurements.

**CCS Concepts:** • Software and its engineering → Software maintenance tools; Software performance; Garbage collection.

**Keywords:** profiling, garbage collection, finalization

### ACM Reference Format:

Sebastian Jordan Montaña, Guillermo Polito, Stephane Ducasse, and Pablo Tesone. 2024. Evaluating Finalization-Based Object Lifetime Profiling. In *Proceedings of the 2024 ACM SIGPLAN International Symposium on Memory Management (ISMM '24)*, June 25, 2024, Copenhagen, Denmark. ACM, New York, NY, USA, 13 pages. <https://doi.org/10.1145/3652024.3665514>

### 1 Introduction

Object lifetime information is crucial to optimize performance through memory optimizations such as pre-tenuring or tuning garbage collector parameters [9, 20]. Implementing algorithms that compute object lifetimes in a precise and scalable manner is nontrivial and requires often a modified virtual machine for execution [4, 6, 7, 14, 15, 32]. For example, Hertz et al. [14, 15] introduced a perfect tracing algorithm that computes object lifetimes called Merlin. However, Merlin has an overhead between 70 and 300 ×.

One practical alternative used in the past is to use finalization mechanisms such as weak references to estimate object lifetimes [1, 25]. However, the topic in question was never explored in depth. In this paper, we explore the profiling of object lifetimes using the ephemeron finalization mechanism [13]. In a nutshell, we instrument object allocations to trace object *birthtime* and we attach an ephemeron to each object to be notified when the object becomes collectible. The main challenge is that naively using such a mechanism

- All the research questions
- More detailed evaluation and methodology
- Detailed results
- Implementation details

# Evaluating Finalization-Based Object Lifetime Profiling

- Finalization profilers can be **weakly actionable** and have **low overhead**
- **Pre-tenuring** ephemerals and **sampling** significantly reduces overhead
- **Sampling** provides **relevant object lifetime information**

Sebastian JORDAN MONTAÑO  
**sebastian.jordan@inria.fr**



[github.com/jordanmontt/illimani-memory-profiler](https://github.com/jordanmontt/illimani-memory-profiler)



*Inria*



Université  
de Lille

23

centralelille  
ÉCOLE CENTRALE DE LILLE

CRISTAL  
Centre de Recherche en Informatique,  
Signal et Automatique de Lille

Région  
Hauts-de-France

